Flooding the filter: a game-theoretic model of AI, submission pressure, and quality collapse

Yuhui Fang

March 18, 2026

Abstract

Generative AI is transforming the nature of academic publishing. By September 2024, more than one-fifth of computer-science preprints showed evidence of large-language-model modification, while monthly submissions to arXiv reached a record 24,226, roughly 50% higher than in late 2022. Verification has not kept pace: reviewer acceptance rates have declined steadily over two decades, and between 6.5% and 16.9% of conference peer reviews are now themselves AI-modified. We study this asymmetry—the fact that AI scales paper production faster than it scales screening—using a game-theoretic model in which AI adoption raises submission volume and lowers drafting cost faster than it raises review capacity, while also increasing the rate at which low-quality manuscripts pass screening. We find three results. First, extreme slot scarcity can stabilize honest submission norms even under degraded screening, because the publication lottery becomes too hard to win. Second, gradual AI adoption can trigger a discrete tipping point into a low-quality long-run outcome; once tipped, the system is self-sustaining, so recovery requires a large coordinated intervention, not incremental reform. Third, expanding publication capacity without scaling verification accelerates the tipping. These results point to a concrete policy agenda of verification-centric institutions (structured audits, submission quotas, and decoupled certification) rather than simply more publication slots.

1 Introduction

We ask what happens to a publishing *system* when it becomes substantially cheaper and faster to produce a manuscript than to review one. By September 2024, more than one-fifth of computer-science preprints on arXiv showed measurable evidence of large-language-model modification (Liang et al., 2025), and at least 13.5% of all biomedical abstracts published that year contained AI-assisted text (Kobak et al., 2025). Monthly submissions to arXiv reached a record 24,226 in October 2024, roughly 50% higher than in late 2022 (arXiv, 2024). Verification has not kept pace. Reviewer acceptance rates have declined steadily over two decades (Meyerson et al., 2025), and between 6.5% and 16.9% of conference peer reviews now contain substantial AI-generated text (Liang et al., 2024). More than half of active researchers surveyed in 2025 reported using AI tools to write reviews, often against explicit journal policy (Naddaf, 2025). The question is not whether individual researchers behave ethically, but whether the system as a whole can absorb an asymmetric shock to production and verification.

While the conventional response to AI in science focuses on disclosure requirements, detection tools, and editorial guidelines, these measures address individual conduct and leave untouched a structural problem: even if every author is honest, reviewers can scrutinize only a finite number of manuscripts carefully per cycle. When submission volume rises faster than review capacity, screening degrades—not because of misconduct, but because of congestion. A reviewer who faces twice as many papers in the same time cannot give each submission the same attention. Congested screening lowers the effective quality threshold, and a lower threshold raises the expected payoff from submitting weak or low-quality work. We focus on this feedback between volume and screening quality as the central mechanism.

We distinguish two bottlenecks in the publishing system, and we argue that AI affects them asymmetrically. The first is verification capacity: the number of papers that reviewers can evaluate carefully in a given period, denoted K_{rev}. The second is allocation capacity: the number of papers the journal can publish per period, denoted K. AI raises submission volume V by reducing the marginal cost of drafting

and resubmitting papers, lowering the effort cost c of each submission. It may also increase the false-positive rate λ—the probability that a low-quality manuscript passes peer review—because AI-assisted writing produces polished text that is harder to distinguish from genuine contributions on surface features alone. Review capacity, by contrast, does not scale symmetrically. Deep expertise, data and code verification, and replication checking remain human-intensive tasks, and journal policies in many fields explicitly restrict AI use in reviews (Naddaf, 2025). Production, in short, scales faster with AI adoption than verification does.

We model this asymmetry in a game-theoretic framework, drawing on strategic models of scientific communication (Zollman et al., 2024; Thurner and Hanel, 2011). We consider a population of authors who choose between two strategies: submitting only genuine work (the cooperative strategy) or submitting everything including low-quality manuscripts (the defecting strategy). A journal applies peer review, subject to both a review capacity K_{rev} and a publication slot constraint K. We introduce an AI adoption intensity parameter $a \in [0, 1]$. Higher a raises submission volume V with elasticity γ, raises review capacity K_{rev} with elasticity $\delta < \gamma$, lowers submission cost c, and raises the false-positive rate λ. The production elasticity exceeds the verification elasticity by assumption, formalizing the asymmetry observed empirically. The net effect of a on the profitability of submitting a low-quality paper, $\Pi_{bad}(x, a) = r \cdot \bar{\rho}(x, a) \cdot \pi_0(x, a) - c(a)$, is positive: AI makes spamming more attractive.

We find three results. First, extreme publication scarcity stabilizes honest submission norms even when screening is imperfect. When the publication lottery is rare enough to win, the expected payoff from submitting a low-quality paper falls below the submission cost, and defection is unprofitable. Scarcity, paradoxically, acts as a deterrent. Second, as a rises, the system can cross a tipping threshold: the long-run outcome shifts abruptly from a high-quality equilibrium to a low-quality one. This tipping occurs because AI simultaneously lowers the break-even point for spam and erodes the screening mechanism that would otherwise deter it. Third, the system exhibits persistent path dependence. The value of a at which quality collapses is strictly lower than the value at which quality recovers. Once the system has tipped into the low-quality long-run outcome, incremental reforms cannot restore it; a large coordinated intervention is needed to cross the recovery threshold.

While expanding publication capacity (raising K) seems like the obvious remedy, we find that doing so without raising review capacity makes the situation worse, not better. More slots increase the expected reward from submitting marginal work, accelerating the drift toward the defecting strategy. Effective intervention must instead target the two primary levers: raising genuine review capacity K_{rev} and reducing the false-positive rate λ. Practical instruments that move these levers include structured code and data audit requirements, submission quotas or tokens, cascading triage systems, and layered certification schemes that decouple preprint dissemination from editorial endorsement. These are not incremental reforms; they constitute a different institutional design, one oriented toward verification rather than volume.

We proceed as follows: Section 4 presents the game-theoretic model; Section 5 derives the replicator dynamics and characterizes the equilibrium regimes; Section 6 examines parameter sensitivity, hysteresis, and strategic levers; and Section 8 concludes.

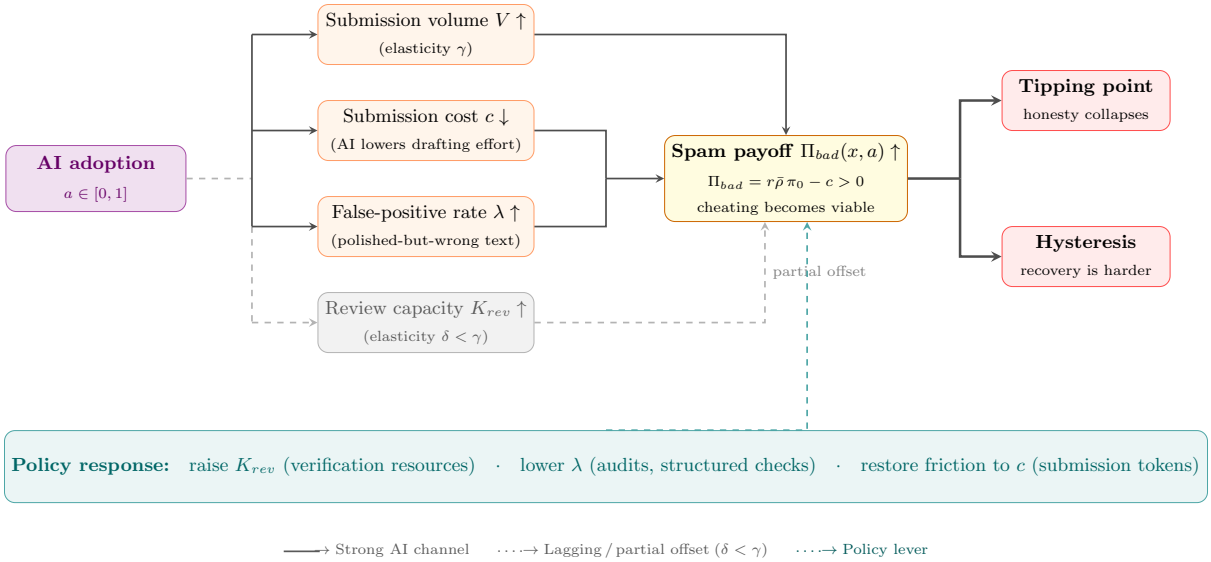

Figure 1: **Mechanism diagram: AI as an asymmetric shock to the publication system.** AI adoption (a) raises submission volume (V, elasticity γ) and lowers drafting cost (c), while polished AI text increases the false-positive rate (λ). Review capacity (K_{rev}) rises more slowly (elasticity $\delta < \gamma$), so the partial offset from AI-assisted reviewing is insufficient to prevent congestion. Together, these shifts raise the profitability of submitting low-quality work ($\Pi_{bad} = r\bar{\rho}\,\pi_0 - c$), pushing the system toward a tipping point and widening the gap between the collapse and recovery thresholds (hysteresis). Effective policy must raise K_{rev}, reduce λ, and restore friction to c (dashed teal arrow).

2 Stylized Facts: AI as an Asymmetric Shock

We draw on evidence that AI dramatically lowers the cost of producing a manuscript. Large language models have cut the time required to draft, revise, and resubmit academic text by a documented margin. In a randomised experiment with 453 college-educated professionals, Noy and Zhang (2023) found that access to ChatGPT reduced the time to complete an incentivised writing task by approximately 40% while simultaneously raising assessed output quality by 18 percentage points, with the largest gains accruing to weaker writers. Stokel-Walker and Van Noorden (2023) report that researchers can revise a full manuscript in under five minutes at an API cost below \$0.50. These micro-level reductions compound at the submission level: iterative resubmission across journals was once a significant labour cost, and AI converts that friction into a near-zero marginal activity, giving every author—regardless of the underlying quality of their work—a cost-effective reason to submit or resubmit. In our model, this channel corresponds to a downward shift in the submission cost parameter $c(a)$.

Aggregate data confirm a sharp rise in submission pressure. Macro-level data are consistent with the micro-level cost story. Monthly submissions to arXiv reached a record 24,226 in October 2024, roughly 50% above the late-2022 baseline; the first month to cross 20,000 submissions was May 2023, shortly after the public release of ChatGPT (arXiv, 2024). More than one-fifth of computer-science preprints sampled in September 2024 showed measurable evidence of large-language-model modification (Liang et al., 2025), and a vocabulary-based analysis of 15 million biomedical abstracts found that at least 13.5% of all 2024 *PubMed* papers contain AI-assisted text, with rates exceeding 40% in some sub-disciplines (Kobak et al., 2025). Alongside this volume growth, annual retraction counts exceeded 10,000 for the first time in 2023, with paper-mill activity implicated in mass retraction events at major publishers. These data show that submission volume V has shifted materially upward as AI adoption has spread, independently of any change in the fraction of authors who genuinely produce high-quality work. In our model, this corresponds to an upward shift in $V(x, a)$ with elasticity γ with respect to a.

While text production scales with AI, verification does not. Rigorous peer review requires reading carefully for logical consistency, checking data against claimed analyses, running or inspecting code, and evaluating genuine novelty—tasks that remain human-intensive in ways that text generation is not. Confidentiality requirements for submitted manuscripts prevent reviewers from uploading papers to third-party AI services, and many journals explicitly forbid AI use in the review process (Naddaf, 2025). Even where AI-assisted reviewing is nominally permitted, the marginal reviewer who can detect a subtle methodological flaw is a scarce specialist whose attention is not fungible with a language model's output. Reviewer acceptance rates have declined continuously for two decades (Meyerson et al., 2025), and the existing burden is already extreme: reviewers contribute more than 100 million hours annually to the peer review system worldwide, with the load concentrated among a small fraction of active researchers (Kovanis et al., 2016). We therefore expect that a system already under strain before the AI era cannot absorb a near-doubling of submission volume through a commensurate increase in review capacity. In our model, this channel is captured by $K_{rev}(a) = K_{rev}(0) \cdot (1 + \delta a)$, where the elasticity δ is strictly smaller than the production elasticity γ, reflecting the structural and policy constraints described above.

AI also polishes manuscripts without improving their reliability, which may raise the false-positive rate. A third channel operates through the quality of the screening signal rather than through volume. AI writing tools produce fluent, well-formatted, and stylistically consistent text (Stokel-Walker and Van Noorden, 2023), and this surface polish reduces the ability of reviewers to distinguish a low-quality submission from a genuine contribution on textual features alone, raising the probability that a flawed manuscript passes the review filter. The screening process degrades further when reviewers themselves turn to AI: between 6.5% and 16.9% of AI-conference peer reviews now contain substantial AI-generated text (Liang et al., 2024), and an experimental audit found that AI-assisted reviews award systematically higher scores, leading to elevated acceptance rates (Latona et al., 2024). Automated detection offers limited recourse—leading AI detectors exhibit false-negative rates above 70%, and adversarial paraphrasing can reduce recall to near zero (Liang et al., 2023). In our model, this channel corresponds to an upward shift in the false-positive rate $\lambda(a)$: as AI adoption rises, the probability that a low-quality submission passes review increases even holding the underlying composition of the author population constant.

We summarize these channels as a single asymmetry hypothesis. The four channels described above—rising submission volume, falling drafting cost, lagging review capacity, and rising false-positive risk—reduce to a single structural claim:

> *AI raises the elasticity of manuscript production above the elasticity of manuscript verification.*

We formalise this as the **AI asymmetry assumption**: with $a \in [0, 1]$ denoting AI adoption intensity,

$$\gamma > \delta \geq 0, \tag{1}$$

where γ is the elasticity of submission volume with respect to a and δ is the elasticity of review capacity with respect to a. This is not an axiom but a testable empirical statement about relative scaling rates. Table 1 translates each channel into a specific model parameter and states the direction of the shift. In Section 6, we examine how our qualitative results depend on the magnitude of the gap $\gamma - \delta$ and on the slope of $\lambda(a)$.

Table 1: **AI adoption channels and their model counterparts.** Each row maps an empirical mechanism to the corresponding model parameter, the direction of the shift, and the evidence base.

Empirical channel	Mechanism	Parameter	Shift	Key evidence
Lower drafting cost	AI reduces time and effort per manuscript; resubmission friction falls	$c(a)$	$\downarrow$	Noy and Zhang (2023); Stokel-Walker and Van Noorden (2023)
Higher submission volume	More authors submit more frequently; low-quality papers enter the pool	V (elasticity γ)	$\uparrow$	arXiv (2024); Liang et al. (2025); Kobak et al. (2025)
Lagging review capacity	Deep verification remains human-intensive; policy restricts AI in review; reviewer shortage pre-dates AI	K_{rev} (elasticity $\delta < \gamma$)	$\uparrow$ slowly	Meyerson et al. (2025); Kovanis et al. (2016); Naddaf (2025)
Rising false-positive rate	Polished AI text harder to distinguish; AI-assisted reviews inflate acceptance scores; detectors unreliable	$\lambda(a)$	$\uparrow$	Liang et al. (2024); Latona et al. (2024); Liang et al. (2023)
Possible quality uplift	AI may also help authors improve genuine contributions (partial offset)	α	$\uparrow$ (offset)	Noy and Zhang (2023)

3 Related Literature

We draw on research showing that peer review functions as a shared resource under increasing strain. Each submission consumes a fixed amount of reviewer time and attention, and when submission volume grows faster than the pool of willing reviewers, that resource degrades. Gross and Bergstrom (2019) show that scientific contests become inefficient at high participation rates: total effort rises, but expected quality per participant falls. Cotton (2013) reaches a parallel conclusion for journals: a small entry cost can deter low-quality submissions and relieve pressure on reviewers. Kovanis et al. (2016) estimate that the global peer review system already consumed over 100 million hours of researcher time each year in the mid-2010s, and Aczél et al. (2021) value the US contribution alone at more than 1.5 billion USD in 2020. Bianchi et al. (2018) model these pressures and find that screening quality declines even when reviewers act in good faith, simply because competing demands leave too little time per paper. The reviewing system, in short, was already near its limits before generative artificial intelligence entered the picture.

While congestion models treat reviewers as passive, a related body of work models authors as strategic agents—and this is where the feedback loop becomes visible. Thurner and Hanel (2011) show that when the review filter is imprecise, submitting everything regardless of quality can become individually rational, pushing the long-run outcome toward low average quality. Zollman et al. (2024) extend this to a multi-journal setting and find that prestige-maximising incentives can reward strategic over honest behaviour. Tiokhin et al. (2021) study honest signalling and show that the good outcome—where only strong work is submitted—is fragile: it collapses when the cost of mimicking quality falls even slightly. Smaldino and McElreath (2016) document this logic empirically, arguing that "publish or perish" incentives act as a selection force that spreads high-volume, low-quality strategies through academic communities. At the aggregate level, Ioannidis (2005) shows that most published findings fail to replicate, not because of fraud, but because structural incentives reward output over reproducibility. Casadevall and Fang (2014) and Stephan (2012) trace this to career and grant structures that treat paper counts as a proxy for scientific value. Paper mills—commercial services producing fabricated manuscripts to order—represent

the extreme end of this logic: annual retraction counts have reached record highs and single incidents have implicated thousands of papers at major publishers. Our model formalises the mechanism underlying these trends: when the review filter is weak and slots are available, submitting low-quality work yields a positive expected payoff, and selection pressure then favours volume over quality.

In contrast to the strategic submission literature, recent empirical work documents how generative artificial intelligence reshapes both sides of the system simultaneously. On the production side, Noy and Zhang (2023) find that access to a large language model cuts professional writing time by 40% and raises rated output quality by 18 percentage points. Liang et al. (2025) detect language-model modification in more than one-fifth of computer-science preprints by late 2024, and Kobak et al. (2025) find AI-assisted text in at least 13.5% of biomedical abstracts from the same period. On the evaluation side, however, the picture is more troubling. Checco et al. (2021) conclude that deep verification—checking data, code, and logical consistency—remains beyond current automated systems. Liang et al. (2024) find that between 6.5% and 16.9% of reviews at major conferences already contain AI-generated text, and Latona et al. (2024) show experimentally that AI-assisted reviews award higher scores, raising acceptance rates. Naddaf (2025) reports that more than half of surveyed researchers now use AI tools for reviewing, often against journal policy. These findings supply the empirical foundation for our asymmetry assumption: artificial intelligence raises submission productivity faster than it raises verification capacity.

A fourth strand addresses the economics of gatekeeping and institutional design. Ellison (2002) shows that the equilibrium publication standard depends on how strongly readers value quality signals. Muller-Itten (2022) finds that journals maintain strict standards to protect their certification value: accepting too freely devalues their own signal. Casadevall and Fang (2014) argue that effective reform must target the false-positive rate directly—through data deposition, structured reporting, and replication requirements—rather than adjusting slot counts alone. Ginsparg (2011) shows that dissemination and certification can be separated: preprint servers make manuscripts available immediately while editorial endorsement follows independently. Cotton (2013) provides the sharpest economic case for submission fees: a small cost of entry restores friction, filters marginal submissions, and reduces reviewer overload.

What we contribute

While these literatures each address part of the problem, none unifies them. Models of strategic submission do not treat artificial intelligence as a structural shock that shifts multiple inputs at once. The empirical AI-in-science literature documents faster production, slower review growth, and higher false-positive risk, but does not connect these to a model with clear predictions about the long-run state of the system. Neither strand addresses the possibility that submission pressure can grow gradually until it crosses a threshold, after which the system shifts suddenly to a low-quality state and resists recovery. The institutional design literature does not ask how its proposals must change when drafting cost approaches zero and the false-positive rate rises continuously with adoption. We close these gaps by connecting all four channels—$c(a)$ falling, V rising, K_{rev} lagging, and $\lambda(a)$ rising—to the long-run state of the publishing system, characterising conditions under which quality collapses suddenly and is hard to restore, and deriving which institutional levers can prevent or reverse that shift.

4 Model

4.1 Setting

We consider a publication game involving a finite population of N authors and a single academic journal. Our framework aligns with established game-theoretic models of scientific publishing Thurner and Hanel (2011), while specifically building on the incentive structures analyzed by Zollman et al. Zollman et al. (2024). Unlike previous static models, we extend the analysis to account for science volume. The fundamental unit of the game is a research paper, which possesses a binary intrinsic quality $q \in \{0, 1\}$. A paper is defined as "good" ($q = 1$) or "bad" ($q = 0$). In each period, authors make submission decisions, and the journal enforces quality control standards subject to a physical capacity constraint.

The publication process unfolds in a strictly sequential timeline, as illustrated in Figure 2:

1. **Submission (Generation of V):** Based on their strategies, authors decide whether to submit their work. The aggregate collection of all submitted manuscripts forms the *Submission Pool*, denoted by $V \leq N$.

2. **Peer Review (Generation of M):** The journal applies a peer review process to filter the submission pool. Acknowledging the inherent unreliability of assessment documented in empirical studies Bornmann (2011), we model reviewers as imperfect classifiers. They are subject to Type I errors (False Positives, rate λ) and Type II errors (False Negatives, rate ϵ). Papers that successfully pass this screening form the set of *Peer-accepted papers*, denoted by M. It is important to note that membership in M signifies that a paper is "scientifically acceptable," but not necessarily "published."

3. **Capacity Check and Storage (K):** The journal has a fixed resource limit or *Journal Storage*, denoted by K. This constraint reflects the physical limits on publication slots or the administrative burden of processing manuscripts Kovanis et al. (2016).

4. **Final Acceptance (ρ):** The final publication decision depends on the relationship between the volume of peer-accepted papers (M) and the journal storage (K). We define ρ as the *probability of being accepted conditional on being peer-accepted*. This variable captures the congestion effect:

$$\rho = \begin{cases} 1 & \text{if } M \leq K \quad \text{(No Congestion)} \\ K/M & \text{if } M > K \quad \text{(Rationing triggered)} \end{cases} \tag{2}$$

In this framework, a paper is published if and only if it passes two hurdles: the quality hurdle (peer review) and the congestion hurdle (storage constraint).

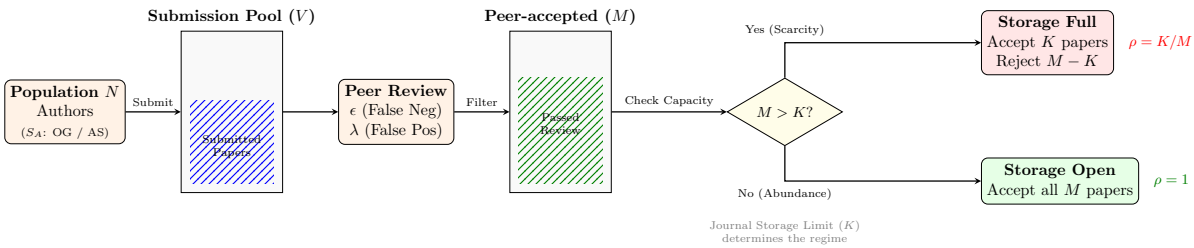

Figure 2: **Schematic representation of the Threshold Competition Model.** The process unfolds in stages: (1) Authors from population N submit papers into pool V; (2) Peer review filters submissions into a set of peer-accepted papers M; (3) The volume M is compared against the journal storage capacity K. If $M > K$, a rationing mechanism is triggered, reducing the conditional acceptance probability ρ below 1.

4.2 Players, strategies, and quality generation

The publication game involves a single journal (J) and a finite population of authors $\mathcal{A} = \{1, 2, \ldots, N\}$. We treat N as a finite exogenous parameter. This finite-population approach allows us to capture the congestion dynamics and "tragedy of the commons" effects inherent in resource-constrained peer review systems, as analyzed in the economic literature on academic publishing (Cotton, 2013; Azar, 2005). Throughout, we use V and M to denote the realised *counts* of submitted and peer-accepted papers, respectively.

Quality generation. In each round, each author generates a manuscript with binary intrinsic quality $q \in \{0, 1\}$, where $q = 1$ denotes a "good" paper and $q = 0$ a "bad" paper. Quality is generated i.i.d. across authors with

$$\Pr(q = 1) = \alpha, \qquad \Pr(q = 0) = 1 - \alpha, \tag{3}$$

where $\alpha \in (0, 1)$ is an exogenous productivity parameter reflecting the baseline scientific capability of the community (Ellison, 2002).

Author strategies (S_A). Each author chooses a submission rule $s_A \in S_A$ that conditions on realised quality. Consistent with signaling models in scientific communication (Tiokhin et al., 2021), we consider four pure strategies:

- **Only Good (OG):** submit iff $q = 1$;
- **Always Submit (AS):** submit regardless of q (strategic spamming);
- **Only Bad (OB):** submit iff $q = 0$ (included for completeness);
- **Never Submit (NS):** never submit.

Equivalently, s_A can be represented by (p_g, p_b), the submission probabilities for good and bad papers:

$$(p_g, p_b) = \begin{cases} (1,0) & \text{if } s_A = \text{OG} \\ (1,1) & \text{if } s_A = \text{AS} \\ (0,1) & \text{if } s_A = \text{OB} \\ (0,0) & \text{if } s_A = \text{NS}. \end{cases} \tag{4}$$

Journal strategies (S_J). The journal commits to an editorial policy $s_J \in S_J$ that maps the perceived quality signal $\hat{q} \in \{0, 1\}$ into an accept/reject decision:

- **Only Good (OG):** accept iff $\hat{q} = 1$;
- **Always Accept (AA):** accept all submissions (ignores $\hat{q}$);
- **Only Bad (OB):** accept iff $\hat{q} = 0$;
- **Reject All (AR):** reject all submissions (ignores $\hat{q}$).

Exogenous parameters and endogenous variables. To clarify the mechanics of congestion, we distinguish between fixed environmental constraints and emergent outcomes:

- **Exogenous parameters:** Population size (N), productivity (α), journal size (K), review noise parameters (ϵ, λ), and payoff primitives, discussed below: c, cost of submission for authors; r, the reward for an author derived from successful publication; B, the benefit to the journal derived from publishing good papers; and D, the the penalty for journals from publishing bad papers.
- **Endogenous variables:** Submission volume (V), peer-accepted volume (M), the rationing factor ($\rho = \max\left\{\frac{K}{M}, 1\right\}$), and the population state x (strategy shares) governing the equilibrium distribution.

4.3 Review noise and acceptance probabilities

Peer review acts as a stochastic filter between the submission volume V and the peer-accepted volume M. Throughout, *pass* refers to passing the review filter and entering the peer-accepted pool M (not final publication, which is also affected by the journal size.

Noise parameters (ϵ, λ). Consistent with empirical meta-analyses on review reliability (Bornmann et al., 2010; Cicchetti, 1991), we assume reviewers are imperfect classifiers. Let $\hat{q} \in \{0, 1\}$ denote the perceived quality signal for a paper with true quality $q \in \{0, 1\}$. We define:

$$\epsilon = \Pr(\hat{q} = 0 \mid q = 1), \qquad \lambda = \Pr(\hat{q} = 1 \mid q = 0), \tag{5}$$

where ϵ is the false negative rate (acting as a tax on honesty) and λ the false positive rate (acting as an incentive for opportunism), a structure supported by empirical studies on scientific gatekeeping (Siler et al., 2015).

Conditional acceptance probabilities. Given a journal policy s_J, define the probability of acceptance conditional on true quality as:

$$\pi_q(s_J) \equiv \Pr(\text{pass} \mid q, s_J), \qquad q \in \{0, 1\}. \tag{6}$$

Under the journal strategies defined in Section 4.2, and aligned with theoretical models of gatekeeper selection (Muller-Itten, 2022), the acceptance vector (π_1, π_0) is given by:

$$(\pi_1(s_J), \pi_0(s_J)) = \begin{cases} (1-\epsilon,\ \lambda) & \text{if } s_J = \text{OG} \\ (1,\ 1) & \text{if } s_J = \text{AA} \\ (\epsilon,\ 1-\lambda) & \text{if } s_J = \text{OB} \\ (0,\ 0) & \text{if } s_J = \text{AR}. \end{cases} \tag{7}$$

Per-author pass-review probability ($\mathcal{P}$). Let $(p_g(s_A), p_b(s_A))$ denote the submission probabilities conditional on realised quality. We define $\mathcal{P}(s_A, s_J)$ as the *ex-ante* probability that an author playing s_A contributes a paper to the peer-accepted pool M under journal policy s_J:

$$\mathcal{P}(s_A, s_J) = \alpha\, p_g(s_A)\, \pi_1(s_J) + (1-\alpha)\, p_b(s_A)\, \pi_0(s_J). \tag{8}$$

Below we enumerate the effective acceptance probabilities for all author strategies under each journal regime.

Case 1: Standard Review (Journal plays Only Good). This is the baseline scenario where $\pi_1 = 1 - \epsilon$ and $\pi_0 = \lambda$.

$$\mathcal{P}(\text{OG}, \text{OG}) = \alpha(1-\epsilon), \tag{9a}$$
$$\mathcal{P}(\text{AS}, \text{OG}) = \alpha(1-\epsilon) + (1-\alpha)\lambda, \tag{9b}$$
$$\mathcal{P}(\text{OB}, \text{OG}) = (1-\alpha)\lambda, \tag{9c}$$
$$\mathcal{P}(\text{NS}, \text{OG}) = 0. \tag{9d}$$

Case 2: Permissive (Journal plays Always Accept). The screening is ignored ($\pi_1 = 1, \pi_0 = 1$), making acceptance depend solely on submission.

$$\mathcal{P}(\text{OG}, \text{AA}) = \alpha, \tag{10a}$$
$$\mathcal{P}(\text{AS}, \text{AA}) = \alpha + (1-\alpha) = 1, \tag{10b}$$
$$\mathcal{P}(\text{OB}, \text{AA}) = 1 - \alpha, \tag{10c}$$
$$\mathcal{P}(\text{NS}, \text{AA}) = 0. \tag{10d}$$

Case 3: Inverse Selection (Journal plays Only Bad). The journal accepts papers perceived as bad ($\hat{q} = 0$). Good papers pass only via false negatives ($\pi_1 = \epsilon$), and bad papers pass via true negatives ($\pi_0 = 1 - \lambda$).

$$\mathcal{P}(\text{OG}, \text{OB}) = \alpha\epsilon, \tag{11a}$$
$$\mathcal{P}(\text{AS}, \text{OB}) = \alpha\epsilon + (1-\alpha)(1-\lambda), \tag{11b}$$
$$\mathcal{P}(\text{OB}, \text{OB}) = (1-\alpha)(1-\lambda), \tag{11c}$$
$$\mathcal{P}(\text{NS}, \text{OB}) = 0. \tag{11d}$$

Case 4: Closed (Journal plays Reject All). No papers enter the peer-accepted pool ($\pi_1 = 0, \pi_0 = 0$).

$$\mathcal{P}(s_A, \text{AR}) = 0 \quad \forall s_A \in S_A. \tag{12}$$

Aggregate peer-accepted volume ($E[M]$). To determine the system-wide congestion, we aggregate these individual passing probabilities across the entire population, a method consistent with burden

estimation in peer review (Kovanis et al., 2016). Let $\mathbf{x} = \{x_{s_A} \mid s_A \in S_A\}$ denote the **population state**, where $x_{s_A} \in [0, 1]$ represents the proportion of authors adopting strategy s_A, satisfying $\sum_{s_A \in S_A} x_{s_A} = 1$.

The expected total volume of peer-accepted papers, $E[M]$, is the sum of contributions from all strategy groups:

$$E[M] = \sum_{i=1}^{N} \mathcal{P}(s_{A,i}, s_J) = N \sum_{s_A \in S_A} x_{s_A} \mathcal{P}(s_A, s_J). \tag{13}$$

This aggregate volume $E[M]$ is the critical endogenous variable that interacts with the journal capacity K to determine the rationing probability ρ defined in Eq. (1).

4.4 Capacity constraint and rationing probability (ρ)

A key feature of the model is the separation between *scientific merit* (passing peer review and entering the set of peer-accepted papers) and *publication slots* (capacity-constrained publication). This distinction captures the reality that top-tier journals often reject scientifically sound papers solely due to space limitations (Ellison, 2002; Casadevall and Fang, 2014).

Journal capacity (K). Let $K \in \mathbb{N}^+$ denote the journal's exogenous publication capacity: at most K papers can be published per cycle due to limited pages and editorial resources (Kovanis et al., 2016; Azar, 2005).

Uniform rationing and the rationing probability (ρ). Let M be the realised number of peer-accepted papers in a given round. If $M \leq K$, all peer-accepted papers are published. If $M > K$, the journal applies *uniform rationing*: it selects exactly K papers uniformly at random from the M peer-accepted papers. Hence, the conditional publication probability for a peer-accepted paper is determined by the realised congestion:

$$\rho(M) \equiv \Pr(\text{published} \mid \text{pass}) = \min\left(1, \frac{K}{M}\right) = \begin{cases} 1 & \text{if } M \leq K, \\ \frac{K}{M} & \text{if } M > K. \end{cases} \tag{14}$$

Ex-post vs. ex-ante congestion. Across rounds, the peer-accepted count M is stochastic because it aggregates individual submission and review outcomes. Authors therefore form expectations over congestion given the population state $\mathbf{x}$ and journal policy s_J, implying an expected rationing factor:

$$\mathbb{E}[\rho] = \mathbb{E}\left[\min\left(1, \frac{K}{M}\right)\right]. \tag{15}$$

Crucially, increases in low-quality or opportunistic submissions raise the expected volume $\mathbb{E}[M]$ and mechanically reduce $\mathbb{E}[\rho]$, generating a negative congestion externality on all authors. This mirrors the inefficiencies found in scientific funding competitions, where excessive participation degrades the system's overall value (Cotton, 2013; Gross and Bergstrom, 2019).

4.5 Payoff analysis

In this subsection, we derive the expected payoffs for both authors and the journal. The introduction of the capacity constraint K fundamentally couples the payoffs: an individual's utility depends not only on their own strategy but also on the aggregate congestion generated by the population. This structure formally aligns with the class of *congestion games*, where individual payoffs are negatively correlated with the total utilization of a shared resource (Rosenthal, 1973).

Notation and general framework Table 2 summarizes the key parameters and variables. We employ a *mean-field approximation* where players base their decisions on the **expected** state of the system. This approach is standard in evolutionary game theory for modeling interactions in large finite populations, allowing deterministic analysis of stochastic processes (Sandholm, 2010; Weibull, 1995).

Table 2: Summary of key parameters and variables used in the payoff analysis.

Symbol	Description
Exogenous Parameters (Fixed)	
N	Total number of authors
K	Journal capacity (maximum publication slots per cycle)
α	Probability that an author produces a good paper
ϵ	False negative rate (Prob. of rejecting a good paper)
λ	False positive rate (Prob. of accepting a bad paper)
r	Reward for a published paper (Author)
c	Cost of submission (Author)
B	Benefit for publishing a good paper (Journal)
D	Penalty for publishing a bad paper (Journal)
C_J	Unit cost of peer review (per submission)
Endogenous Variables (Dependent on Strategy)	
V	Total number of submissions (Submission Pool)
M	Number of papers passing peer review (Perception Pool)
Y	Number of papers finally published ($Y \leq K$)
$\bar{\rho}$	Expected rationing probability ($\mathbb{E}[\min(1, K/M)]$)
x	Proportion of authors adopting a specific strategy (Population State)

Submission and publication probabilities. For an author playing strategy $s_A \in S_A$, we define the unconditional probability of submitting a paper as $P_{\text{sub}}(s_A)$. Based on the submission vectors (p_g, p_b) defined in Section 4.2, this is:

$$P_{\text{sub}}(s_A) = \alpha p_g(s_A) + (1-\alpha) p_b(s_A). \tag{16}$$

Recall from Section 4.3 that $\mathcal{P}(s_A, s_J)$ denotes the probability of *entering the peer-accepted pool* M. However, entering M does not guarantee publication. Let $\bar{\rho}$ denote the **expected rationing probability**:

$$\bar{\rho} \equiv \mathbb{E}[\rho] = \mathbb{E}\left[\min\left(1, \frac{K}{M}\right)\right]. \tag{17}$$

The joint probability of a paper being finally published is the product of passing review and passing the capacity check: $\bar{\rho} \cdot \mathcal{P}(s_A, s_J)$.

Volumes: Submitted (V), Passed (M), and Published (Y). It is crucial to distinguish between papers that pass peer review (M) and those that are finally published (Y). Let V_G and V_B be the expected volumes of good and bad submissions generated by the population. Let M_G and M_B be the expected volumes of papers entering the peer-accepted pool. We define Y_G and Y_B as the expected volumes of **finally published** papers:

$$Y_G = \bar{\rho} M_G, \quad Y_B = \bar{\rho} M_B. \tag{18}$$

General payoff functions

Author's utility (U_A). Let r be the reward for a published paper and c be the cost of submission. The expected payoff for an author playing s_A against a journal s_J is the expected reward from publication

minus the reliable cost of submission:

$$U_A(s_A, s_J) = r \cdot \underbrace{\bar{\rho} \cdot \mathcal{P}(s_A, s_J)}_{\text{Prob(Published)}} - c \cdot \underbrace{P_{\text{sub}}(s_A)}_{\text{Prob(Submitted)}} \tag{19}$$

Journal's utility (U_J). Let B be the benefit of publishing a good paper, D be the penalty for a bad paper, and C_J be the unit cost of reviewing a submission. The journal's utility depends on the *published* volumes (Y_G, Y_B) and the *submitted* volume $(V = V_G + V_B)$. This formulation captures the trade-off between quality assurance and processing costs (Muller-Itten, 2022):

$$U_J(s_J) = B \cdot Y_G - D \cdot Y_B - C_J \cdot V \tag{20}$$

Substituting $Y = \bar{\rho}M$, we obtain the working form of the journal's payoff:

$$U_J(s_J) = \bar{\rho}\,[B \cdot M_G(s_J) - D \cdot M_B(s_J)] - C_J \cdot V \tag{21}$$

Remark: Eq. (21) relies on the approximation $\mathbb{E}[\rho M] \approx \mathbb{E}[\rho]\mathbb{E}[M]$. As N grows large, the relative fluctuations in M diminish, making this mean-field limit exact (Sandholm, 2010).

Derivation of strategy-specific payoffs We now explicitly list the payoff functions for the 8 primary scenarios (4 Author strategies and 4 Journal strategies).

Author payoffs (under Journal Strategy s_J = OG). In the standard academic context, the journal plays **Only Good (OG)**. The review filter applies $(\pi_1, \pi_0) = (1 - \epsilon, \lambda)$. Substituting the specific $\mathcal{P}$ and P_{sub} values into Eq. (19), we derive:

$$U_A(\text{OG}, \text{OG}) = \alpha\,[r\bar{\rho}(1 - \epsilon) - c] \tag{22}$$

$$U_A(\text{AS}, \text{OG}) = r\bar{\rho}\,[\alpha(1 - \epsilon) + (1 - \alpha)\lambda] - c \tag{23}$$

$$U_A(\text{OB}, \text{OG}) = (1 - \alpha)\,[r\bar{\rho}\lambda - c] \tag{24}$$

$$U_A(\text{NS}, \text{OG}) = 0 \tag{25}$$

Journal payoffs (under Population State x). The journal chooses a strategy s_J to maximize utility given the incoming flow of submissions. Let $V_G(\mathbf{x})$ and $V_B(\mathbf{x})$ be the volumes of good and bad submissions determined by the population state.

1. Only Good (OG): The journal filters papers. The passed volumes are $M_G = V_G(1 - \epsilon)$ and $M_B = V_B\lambda$.

$$U_J(\text{OG}) = \bar{\rho}\,[BV_G(1 - \epsilon) - DV_B\lambda] - C_J(V_G + V_B) \tag{26}$$

2. Always Accept (AA): The journal accepts everything $(M_G = V_G, M_B = V_B)$. Congestion is maximized here, so $\bar{\rho}$ will be minimal.

$$U_J(\text{AA}) = \bar{\rho}\,[BV_G - DV_B] - C_J(V_G + V_B) \tag{27}$$

3. Only Bad (OB): The journal applies inverse selection $(M_G = V_G\epsilon, M_B = V_B(1 - \lambda))$.

$$U_J(\text{OB}) = \bar{\rho}\,[BV_G\epsilon - DV_B(1 - \lambda)] - C_J(V_G + V_B) \tag{28}$$

4. Reject All (AR): The journal publishes nothing $(Y = 0)$. Review costs are still incurred for processed submissions.

$$U_J(\text{AR}) = -C_J(V_G + V_B) \quad \text{(normalized to 0 in dominance analysis)} \tag{29}$$

Summary of transition probabilities For completeness, Table 3 summarizes the effective probability $\mathcal{P}(s_A, s_J)$ of entering the peer-accepted pool for any strategy pair. These values are the key inputs for the payoff equations.

Table 3: Matrix of Effective Passing Probabilities $\mathcal{P}(s_A, s_J)$. This represents the probability that an author's paper enters the peer-accepted pool M (before rationing).

Author \ Journal	**Only Good (OG)**	**Always Accept (AA)**	**Only Bad (OB)**	**Reject All (AR)**
OG (α good)	$\alpha(1-\epsilon)$	α	$\alpha\epsilon$	0
AS (All)	$\alpha(1-\epsilon)+(1-\alpha)\lambda$	1	$\alpha\epsilon+(1-\alpha)(1-\lambda)$	0
OB ($1-\alpha$ bad)	$(1-\alpha)\lambda$	$1-\alpha$	$(1-\alpha)(1-\lambda)$	0
NS (None)	0	0	0	0

4.6 Dominance analysis and reduced game

Before analyzing the evolutionary dynamics, we simplify the strategy space by iteratively eliminating strictly dominated strategies. This procedure, known as *iterated elimination of strictly dominated strategies* (IESDS), allows us to focus on the rationalizable core of the game (Fudenberg and Tirole, 1991). We conduct this analysis under the standard **mean-field approximation**, holding the aggregate congestion factor $\bar{\rho}$ fixed. The elimination proceeds in two stages: first, we identify strategies that are dominated purely by the structural logic of the game (independent of specific cost-benefit inequalities); second, we introduce parameter constraints to eliminate strategies based on economic trade-offs.

Stage 1: Structural elimination (Unconditional dominance) We first establish dominance relations that hold under weak axiomatic assumptions regarding the nature of peer review and scientific value, consistent with signaling models in scientific communication (Tiokhin et al., 2021).

Elimination of Journal "Only Bad" (OB). The strategy **Only Bad (OB)** represents an inverse selection policy. We compare it to **Only Good (OG)**. Using the payoff functions derived in Section 4.5, the difference in expected utility is:

$$\begin{aligned}\Delta U_J &= U_J(\text{OG}) - U_J(\text{OB}) \\ &= \bar{\rho}\left[BV_G(1-\epsilon) - DV_B\lambda\right] - \bar{\rho}\left[BV_G\epsilon - DV_B(1-\lambda)\right] \\ &= \bar{\rho}\left[BV_G(1-2\epsilon) + DV_B(1-2\lambda)\right] \end{aligned} \tag{30}$$

Condition: We assume peer review is **informative**, defined as performing better than random guessing:

$$\epsilon < 0.5 \quad \text{and} \quad \lambda < 0.5. \tag{31}$$

Under this definition, both $(1-2\epsilon)$ and $(1-2\lambda)$ are strictly positive. Since $B, D, V_G, \bar{\rho} > 0$, it follows that $\Delta U_J > 0$. Thus, **Journal OB is strictly dominated by OG** solely by virtue of the definition of peer review quality, regardless of the magnitude of rewards or penalties.

Elimination of Author "Only Bad" (OB). Consider the choice between **Always Submit (AS)** and **Only Bad (OB)**. AS submits both paper types ($q \in \{0,1\}$), while OB submits only bad papers ($q = 0$). The payoff difference is the marginal utility of submitting a good paper:

$$\begin{aligned}\Delta U_A &= U_A(\text{AS}, \text{OG}) - U_A(\text{OB}, \text{OG}) \\ &= [U_A(q=1) + U_A(q=0)] - U_A(q=0) \\ &= U_A(q=1) = \alpha\left[r\bar{\rho}(1-\epsilon) - c\right] \end{aligned} \tag{32}$$

Condition: We assume the **Axiom of Value**: in a functional scientific system, a correctly identified high-quality paper yields a positive net return (i.e., $r(1-\epsilon) > c$). If this axiom holds, then $\Delta U_A > 0$ for any non-zero congestion $\bar{\rho}$. Thus, **Author OB is strictly dominated by AS** due to monotonicity: a rational author will never withhold a high-quality paper while simultaneously submitting a low-quality one.

Parameter constraints for a non-trivial game Having eliminated the structurally irrational strategies, we now focus on the remaining set. To ensure the game yields interesting evolutionary dynamics (avoiding trivial collapse or utopian perfection), we define the "viable region" of the parameter space using the **abundance benchmark** ($\bar{\rho} = 1$).

1. **Viability Constraint (Honesty is profitable):**

$$r(1 - \epsilon) > c \tag{33}$$

   This formally restates the Axiom of Value. It ensures that without congestion, the honest strategy (OG) dominates silence (NS), a necessary condition for any scientific activity (Gross and Bergstrom, 2019).

2. **Profitability Constraint (Spamming is profitable):**

$$\lambda r > c \tag{34}$$

   This constraint drives the congestion dilemma. It implies that in a frictionless system ($\bar{\rho} = 1$), the expected reward of a bad paper exceeds the submission cost. If $c > \lambda r$, bad papers are never submitted, and the game trivially converges to perfect honesty.

3. **Reputation Constraint (Screening is preferred):**

$$DV_B(1 - \lambda) > BV_G\epsilon \tag{35}$$

   This condition ensures the journal prefers screening (OG) over blind acceptance (AA). It states that the aggregate penalty avoided by rejecting bad papers outweighs the benefit lost from accidentally rejecting good ones (Muller-Itten, 2022).

Stage 2: Conditional elimination (Economic dominance) We now eliminate strategies based on the specific economic trade-offs defined above.

Dominance of Journal "Always Accept" (AA). We compare **Only Good (OG)** with **Always Accept (AA)**. Assuming the processing cost C_J applies to all submissions (Processing Cost Model), the difference depends on the net value of papers rejected by OG but accepted by AA:

$$\begin{aligned} U_J(\text{OG}) - U_J(\text{AA}) &= \bar{\rho}\left[(BM_G^{\text{OG}} - DM_B^{\text{OG}}) - (BM_G^{\text{AA}} - DM_B^{\text{AA}})\right] \\ &= \bar{\rho}\left[D(V_B - M_B^{\text{OG}}) - B(V_G - M_G^{\text{OG}})\right] \\ &= \bar{\rho}\left[DV_B(1 - \lambda) - BV_G\epsilon\right] \end{aligned} \tag{36}$$

Under the **Reputation Constraint** (35), the penalty avoided (first term) exceeds the benefit lost (second term). Thus, **AA is strictly dominated by OG**. This result aligns with gatekeeper theory, suggesting that competitive journals maintain strict standards to maximize their distinctiveness and reputation (Muller-Itten, 2022).

Dominance of Journal "Reject All" (AR). Strategy AR yields a normalized payoff of 0 (or $-C_J V$). Since OG generates positive utility from good papers ($B > 0$) and filters bad ones, OG dominates AR provided the net value of the peer-accepted pool is positive: $BM_G > DM_B$. We assume the journal operates in a parameter region where valid science is net positive.

Dominance of Author "Never Submit" (NS). Comparing **Only Good (OG)** with **Never Submit (NS)** ($U_A = 0$):

$$U_A(\text{OG}, \text{OG}) - U_A(\text{NS}, \text{OG}) = \alpha\left[r\bar{\rho}(1 - \epsilon) - c\right] \tag{37}$$

Under the **Viability Constraint** (33), this difference is positive as long as the congestion does not exceed the critical threshold $\bar{\rho} > \frac{c}{r(1-\epsilon)}$. In the viable regime, **NS is strictly dominated by OG** (and by AS). This condition reflects the rational participation constraint found in scientific contests, where authors only enter if the expected payoff exceeds the effort cost (Gross and Bergstrom, 2019).

The reduced evolutionary game Following the two-stage elimination process, the complex multi-strategy interaction reduces to a concise 2×1 asymmetric game.

1. **The Journal** is a fixed institutional player locked into the **Only Good (OG)** strategy. It acts as the environment, applying a filter (ϵ, λ) and a capacity constraint K.
2. **The Author Population** evolves between two surviving pure strategies:
   - **Only Good (OG) [Cooperators]:** Authors who submit only high-quality work ($q = 1$).
   - **Always Submit (AS) [Defectors]:** Authors who submit everything ($q \in \{0, 1\}$), effectively spamming the journal with bad papers to exploit λ.

We define the **population state** $x(t) \in [0, 1]$ as the fraction of authors adopting the **Only Good** strategy. The evolutionary dynamics of the system are thus entirely determined by the time evolution of x. This reduction captures the essential dynamic of scientific evolution: the selection pressure between high-effort integrity and low-effort quantity (Smaldino and McElreath, 2016).

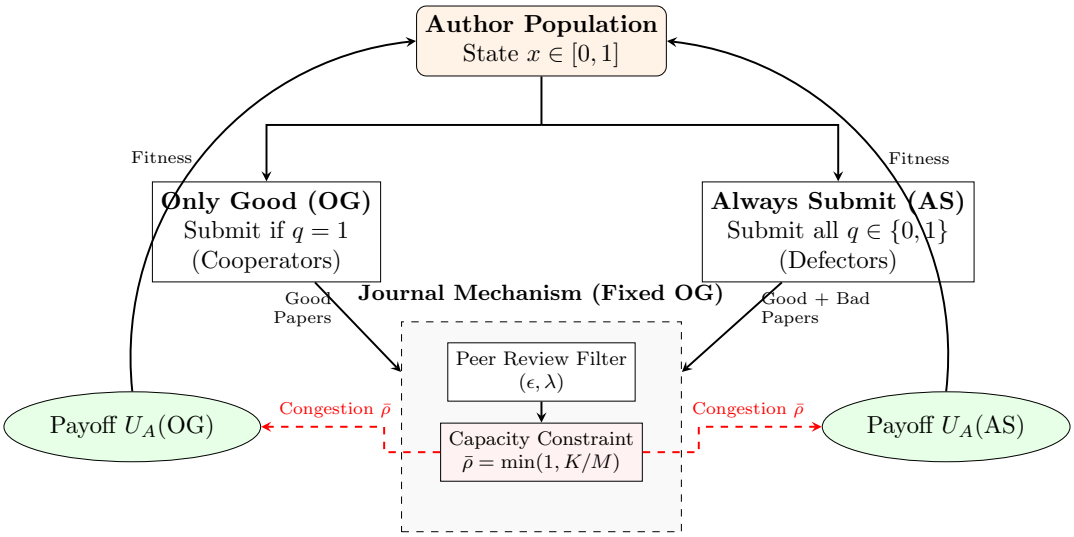

Figure 3: **Structure of the Reduced Evolutionary Game.** The author population evolves between two strategies: Only Good (OG) and Always Submit (AS). The Journal acts as a fixed mechanism consisting of a peer review filter and a capacity constraint. Submissions from AS authors increase the total volume M, triggering the capacity constraint. The resulting **Congestion Externality** ($\bar{\rho}$) reduces the publication probability for all papers, feeding back into the payoffs and driving the evolutionary dynamics.

5 Replicator Dynamics and Bifurcation Analysis

5.1 Population state and congestion functions

Following the dominance analysis, the strategy space reduces to a contest between two author strategies, with the journal playing a fixed filtering role. We define the **population state** $x(t) \in [0, 1]$ as the fraction of authors playing **Only Good (OG)**, while $1 - x(t)$ play **Always Submit (AS)**. We assume the system operates in the region $\lambda r < c < r$: in the absence of congestion, submitting low-quality work is unprofitable, but it becomes viable once the screening mechanism degrades sufficiently—a mechanism analogous to the market for lemons (Akerlof, 1970; Thurner and Hanel, 2011).

To capture the dual impact of congestion, we model two distinct bottlenecks. Numerical simulations use a smoothed approximation to capture pre-congestion friction; the analytical forms below use standard cutoffs for clarity.

Stage 1: The screening intensity function $\eta(x)$ The total submission volume is $V(x) = N[1-(1-\alpha)x]$. The screening intensity $\eta(x)$ represents the probability that a paper undergoes rigorous peer review. It degrades when volume exceeds the review capacity K_{rev}, a phenomenon consistent with models of "reviewer fatigue" and resource scarcity (Azar, 2005; Kovanis et al., 2016):

$$\eta(x) = \min\left(1, \frac{K_{rev}}{V(x)}\right) \tag{38}$$

We assume "rubber-stamping" implies unconditional acceptance of unscreened papers. This reflects the transition from rational selection to random selection observed in overloaded systems (Thurner and Hanel, 2011). Thus, the pass probability for a bad paper, $\pi_0(x)$, increases as screening fails:

$$\pi_0(x) = \eta(x)\lambda + (1-\eta(x)) \cdot 1 = 1 - \eta(x)(1-\lambda) \tag{39}$$

For a good paper, the pass probability $\pi_1(x)$ also approaches 1 as scrutiny relaxes (assuming false negatives ϵ only occur during rigorous review):

$$\pi_1(x) = \eta(x)(1-\epsilon) + (1-\eta(x)) \cdot 1 = 1 - \eta(x)\epsilon \tag{40}$$

Stage 2: The rationing probability function $\bar{\rho}(x)$ Let K be the publication slot capacity. The expected volume of papers entering the "perceived good" pool $M(x)$ is:

$$\begin{aligned} M(x) &= V_G(x)\pi_1(x) + V_B(x)\pi_0(x) \\ &= N\alpha \cdot \pi_1(x) + N(1-\alpha)(1-x) \cdot \pi_0(x) \end{aligned} \tag{41}$$

Using a mean-field approximation, the expected rationing probability is:

$$\bar{\rho}(x) \approx \min\left(1, \frac{K}{M(x)}\right) \tag{42}$$

This formulation captures the "lottery nature" of high-impact publishing: when the number of scientifically sound papers (M) exceeds the journal's physical space (K), acceptance becomes a game of chance rather than pure merit (Casadevall and Fang, 2014).

5.2 Derivation of the replicator equation

Since both OG and AS strategies submit good papers whenever they occur (with identical acceptance probabilities), the expected payoff component from good papers cancels out in the differential. The dynamics are driven entirely by the net utility of the additional "bad" papers submitted by AS authors, denoted as $\Pi_{bad}(x)$. The replicator equation is:

$$\dot{x} = -x(1-x)(1-\alpha)\Pi_{bad}(x) \tag{43}$$

A bad paper yields reward r only if it passes both bottlenecks. The expected net utility is:

$$\Pi_{bad}(x) = r \cdot \bar{\rho}(x) \cdot \pi_0(x) - c \tag{44}$$

5.3 Bifurcation and stability analysis

The system exhibits **Bistability** because the two congestion effects work in opposition, creating a dynamic tension characteristic of tipping-point models (Schelling, 1978):

- **Review Collapse ($\pi_0 \uparrow$):** Increases the incentive to cheat (pro-corruption).
- **Storage Rationing ($\bar{\rho} \downarrow$):** Decreases the incentive to cheat (pro-honesty).

When the review collapse effect dominates at low honesty, $\Pi_{bad}(x)$ becomes positive, creating a *corruption trap* where opportunistic strategies proliferate (Smaldino and McElreath, 2016). Conversely, when strict screening dominates at high honesty, $\Pi_{bad}(x)$ becomes negative, creating a basin of honesty.

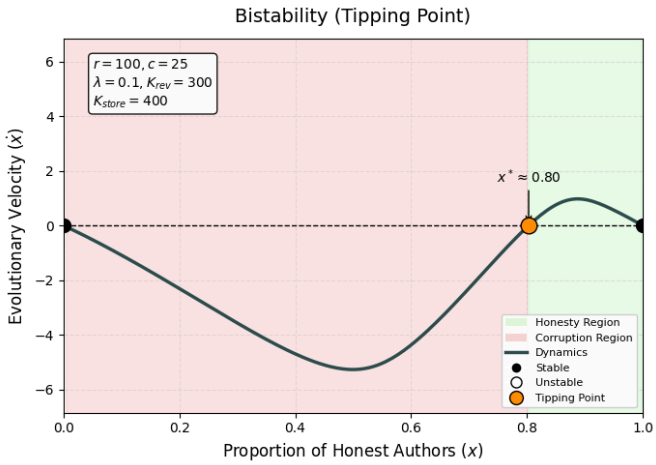

Figure 4: **Phase Portrait of the Dual-Bottleneck Dynamics (Bistability).** Parameters: $N = 1000, K_{rev} = 300, K_{store} = 400, r = 100, c = 25, \lambda = 0.1$. The curve shows the evolutionary velocity $\dot{x}$. The system exhibits three equilibria: (1) **Trap of Corruption** ($x = 0$): Stable ($\Pi_{bad}(0) > 0$). (2) **Basin of Honesty** ($x = 1$): Stable ($\Pi_{bad}(1) < 0$). (3) **Tipping Point** ($x^* \approx 0.82$): An unstable internal repeller. The orange marker indicates the critical threshold separating the two basins of attraction.

Analytical Condition for Instability of the Internal Equilibrium To rigorously characterize the stability of an interior equilibrium $x^* \in (0, 1)$, we analyze the local derivative of the velocity field. Following standard evolutionary stability analysis (Hofbauer and Sigmund, 1998), the equilibrium is unstable (a repeller) if and only if the net utility of cheating is locally decreasing in x:

$$\left.\frac{d\Pi_{bad}}{dx}\right|_{x^*} < 0 \tag{45}$$

Recall that $\Pi_{bad}(x) = r \cdot \bar{\rho}(M(x)) \cdot \pi_0(x) - c$. Using the chain rule, we decompose the total derivative into two competing forces:

$$\frac{d\Pi_{bad}}{dx} = r \left(\underbrace{\bar{\rho}(M)\frac{d\pi_0}{dx}}_{\text{Review Effect}} + \underbrace{\pi_0(x)\frac{d\bar{\rho}}{dM}\frac{dM}{dx}}_{\text{Rationing Effect}} \right) \tag{46}$$

Note that $\frac{d\pi_0}{dx} < 0$ (since higher honesty reduces volume V, improving review quality η, thus lowering π_0). Conversely, the rationing term is typically positive (higher honesty reduces false positives in M, relaxing rationing, thus increasing $\bar{\rho}$). Therefore, the tipping point is unstable if and only if the Review-Dominated Condition holds locally:

$$\underbrace{\left|\bar{\rho}(M)\frac{d\pi_0}{dx}\right|}_{\text{Loss from stricter review}} > \underbrace{\pi_0(x)\frac{d\bar{\rho}}{dM}\frac{dM}{dx}}_{\text{Gain from relaxed rationing}} \tag{47}$$

This condition confirms that bistability arises when the recovery of peer review quality (as x increases) outweighs the relaxation of publication slots.

5.4 Equilibrium regimes

Stable coexistence: rationing-dominated dynamics While the default scenario exhibits bistability, a fundamentally different topological regime emerges when the storage constraint (K_{store}) is significantly tighter than the review constraint (K_{rev}). Figure 5 illustrates this scenario, where the evolutionary velocity $\dot{x}$ transitions from positive to negative, creating a Stable Internal Equilibrium ($x^* \approx 0.44$).

Existence and Stability Conditions For a stable coexistence equilibrium to exist, two conditions must be met:

1. **Boundary Conditions:** The system must be repelled from both extremes. This requires $\Pi_{bad}(1) > 0$ (Temptation at Honesty) and $\Pi_{bad}(0) < 0$ (Unprofitability at Corruption). By the Intermediate Value Theorem, there must exist at least one root $x^* \in (0, 1)$ where $\Pi_{bad}(x^*) = 0$.
2. **Stability Condition:** The slope of the velocity field must be negative, i.e., $\frac{d\Pi_{bad}}{dx}\Big|_{x^*} > 0$. This corresponds to the Rationing-Dominated Condition*, where the "punishment" from overcrowding (tightening $\bar{\rho}$) outweighs the "opportunity" from collapsed screening.

Real-world implication: The "Defense by Congestion" paradox The emergence of a stable internal equilibrium reveals a paradoxical phenomenon: **extreme congestion acts as an unintentional defense mechanism**. In journals with strict page limits ($K_{store} \ll N$), even if peer review fails ($\pi_0 \to 1$), the sheer volume of submissions drives the acceptance rate $\bar{\rho}$ so low that spamming becomes statistically unprofitable. This explains why top-tier journals do not degenerate into full corruption even under heavy load—the "lottery" is simply too hard to win. This dynamic aligns with contest models where excessive competition discourages entry by marginal participants (Gross and Bergstrom, 2019). However, it also implies vulnerability: if capacity is expanded (increasing K_{store}), the system loses this passive defense, potentially shifting the topology to bistability or collapse. This highlights the strategic value of artificial scarcity in maintaining the "gatekeeper" function of academic journals (Muller-Itten, 2022).

Boundary regimes: global honesty and global corruption Beyond the interior dynamics, the system may exhibit monotonic behaviour where the incentive to submit low-quality work is uniformly positive or negative.

Global Honesty ($x^* = 1$): Defense by Rigor When $\Pi_{bad}(x) < 0$ for all $x \in [0, 1]$, the evolutionary velocity $\dot{x}$ remains positive throughout the domain, driving the system unconditionally toward $x^* = 1$ (Figure 6). A sufficient condition for this regime is that the payoff is negative even at the most favorable boundary ($x = 1$):

$$\Pi_{bad}(1) = r \cdot \lambda - c < 0 \tag{48}$$

This implies that suppressing the false-positive rate λ below the cost-reward ratio c/r is the primary safeguard for global honesty.

Global Corruption ($x^* = 0$): Systemic Failure Conversely, when $\Pi_{bad}(x) > 0$ for all $x \in [0, 1]$, the system collapses unconditionally to $x^* = 0$ (Figure 7). This occurs when the payoff remains positive even under maximum congestion ($x = 0$):

$$\Pi_{bad}(0) = r \cdot \bar{\rho}(V_{max}) \cdot \pi_0(V_{max}) - c > 0 \tag{49}$$

This regime represents a total failure of governance, often seen in predatory journals where the profit motive drives acceptance regardless of quality.

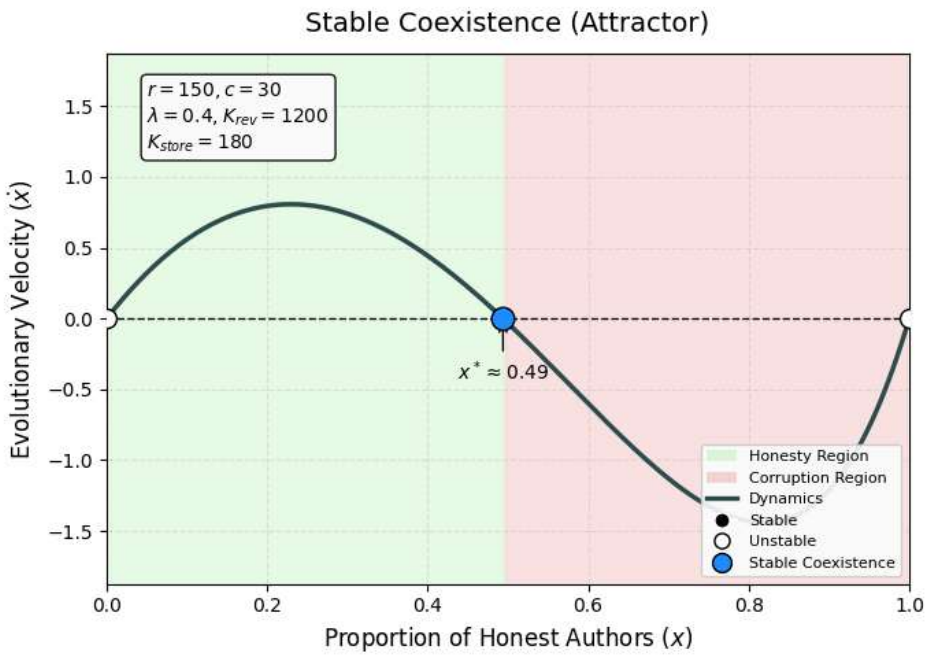

Figure 5: **Phase Portrait of Stable Coexistence (Rationing-Dominated).** Parameters: $N = 1000, K_{rev} = 1000, K_{store} = 180, r = 150, c = 30, \lambda = 0.4$. The system exhibits a stable attractor at $x^* \approx 0.49$. **Left Zone (Self-Correction):** At low honesty ($x < x^*$), severe storage rationing reduces the payoff below zero ($\Pi_{bad} < 0$), forcing the population to become more honest. **Right Zone (Temptation):** At high honesty ($x > x^*$), the relaxed rationing makes cheating profitable again ($\Pi_{bad} > 0$), forcing the population towards corruption. The system settles into a permanent mix of honest and opportunistic authors (Blue Point).

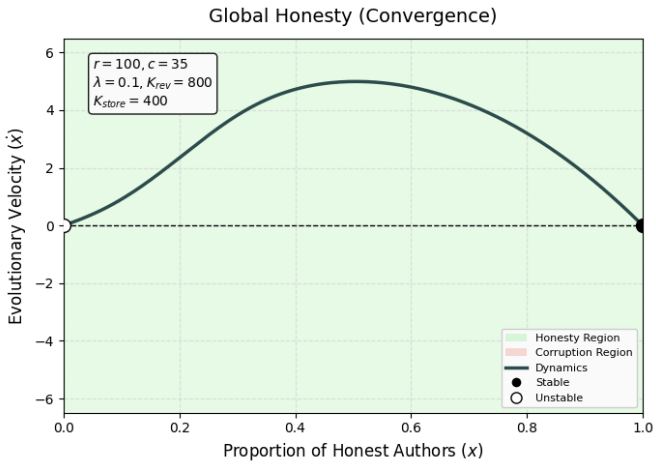

Figure 6: **Phase Portrait of Global Honesty.** Parameters: $N = 1000, K_{rev} = 800, K_{store} = 400, r = 100, c = 35, \lambda = 0.1$. The velocity is strictly positive (Green Zone), ensuring convergence to $x = 1$ from any initial state.

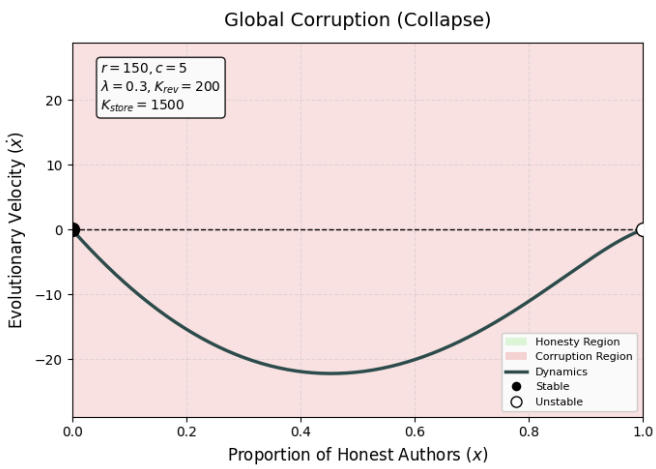

Figure 7: **Phase Portrait of Global Corruption.** Parameters: $N = 1000, K_{rev} = 200, K_{store} = 1500, r = 150, c = 5, \lambda = 0.3$. The velocity is strictly negative (Red Zone), ensuring collapse to $x = 0$.

5.5 Summary: A Topological Classification of Regimes

The dynamical behavior is ultimately determined by the interplay between the boundary payoffs. Table 4 summarizes the four topological regimes.

Table 4: **Topological Classification of Dynamical Regimes based on Boundary Payoffs**

Regime	$\mathbf{\Pi_{bad}(1)}$ *(Low-spam Limit)*	$\mathbf{\Pi_{bad}(0)}$ *(Max-load Limit)*	**Dynamical Topology & Implication**
GH	< 0	< 0	**Honesty-dominant.** $x = 1$ is locally stable; $x = 0$ is unstable. Trajectories converge to $x = 1$ globally, *provided* $\Pi_{bad}(x)$ *has no interior root.*
GC	> 0	> 0	**Corruption-dominant.** $x = 0$ is locally stable; $x = 1$ is unstable. Trajectories collapse to $x = 0$ globally, *provided* $\Pi_{bad}(x)$ *has no interior root.*
BS	< 0	> 0	**Bistability.** Both boundaries are locally stable attractors. By continuity, **at least one** unstable interior equilibrium x^* exists to separate the basins.
SC	> 0	< 0	**Stable Coexistence.** Both boundaries are unstable repellers. By continuity, **at least one** stable interior equilibrium x^* exists (trapped by congestion). Uniqueness depends on monotonicity.

Note: The **Low-spam Limit** ($x = 1$) is driven primarily by review rigor (λ) via $\Pi_{bad}(1) \approx r\lambda - c$. The **Max-load Limit** ($x = 0$) is driven by the collapse of acceptance rate ($\bar{\rho}$) under maximum volume, regardless of review failure.

This classification reveals a dual-control mechanism: **Review Quality (λ)** secures the honest boundary ($x = 1$), while **Storage Capacity (K_{store})** secures the corrupt boundary ($x = 0$).

6 Parameter Sensitivity and Regime Transitions

6.1 Methodology and Baseline Calibration

We want to know which policies maximize publication quality and minimize academic misconduct. However, due to the non-linear interplay between the review bottleneck (η) and the storage bottleneck ($\bar{\rho}$), the system does not respond smoothly to policy changes. Instead, the ecosystem undergoes phase transitions between qualitatively distinct regimes.

To analyze these transitions, we use numerical bifurcation analysis. We define the instantaneous Average Publication Quality $Q(x)$ as the proportion of good papers within the set of accepted publications, a metric that aligns with standard meta-scientific frameworks for assessing the reliability of research findings (Ioannidis, 2005):

$$Q(x) = \frac{N\alpha \cdot \pi_1(x)}{M(x)} = \frac{\alpha\pi_1(x)}{\alpha\pi_1(x) + (1-\alpha)(1-x)\pi_0(x)} \tag{50}$$

Crucially, the long-term quality Q_∞ is determined by the system's stable attractors. In bistable regimes, Q_∞ is path-dependent, meaning the final quality depends on whether the system started from a high-trust or low-trust state.

We classify the system into four regimes based on the stability of equilibria:

1. **Global Honesty (GH):** the system converges to $x = 1$, yielding $Q_\infty = 1.0$.
2. **Global Corruption (GC):** the system collapses to $x = 0$, yielding $Q_\infty \approx \alpha$—a "market for lemons" where the journal fails to filter noise (Akerlof, 1970).
3. **Bistability (BS):** two stable attractors coexist and the outcome depends on history.
4. **Stable Coexistence (SC):** an interior equilibrium x^* is maintained by congestion, stabilizing quality at a mediocre level ($Q_\infty \approx 0.5$).

To ensure comparative consistency, we use two representative parameter sets (Table 5). Set A represents a typical journal with moderate capacity, prone to bistability. Set B represents a top-tier journal with strict page limits and high rewards, prone to stable coexistence.

Table 5: Parameter Calibration for Numerical Analysis

Symbol	Description	Set A (Baseline)	Set B (High-Pressure)
N	Total Population	1000	1000
α	Author Productivity	0.2	0.2
K_{rev}	Review Capacity	**300**	**1200**
K	Slot Capacity	**400**	**180**
r	Reward	100	**150**
c	Cost	25	**30**
λ	False Positive Rate	0.1	**0.4**
ϵ	False Negative Rate	0.05	0.01

Note: Set A is used as the default for sensitivity sweeps. Set B is used to illustrate the defense-by-congestion mechanism.

Why two parameter sets? A single baseline cannot capture the full range of the system's dynamical behaviours, because the dual-bottleneck mechanism is non-monotonic. We distinguish two archetypes. Set A represents the default state of many modern journals: moderate prestige ($r = 100$) and limited but not extreme storage constraints ($K = 400$). Here the primary failure mode is the review bottleneck ($K_{rev} = 300$)—when submission volume rises, the review process collapses before the storage limit is reached. This setup reflects the empirical reality of the "global burden" of peer review, where reviewer scarcity is the dominant constraint (Kovanis et al., 2016), and it generates a bistable topology in which the system is fragile and prone to catastrophic tipping. Set B represents top-tier prestige journals such as *Nature* or *Science*, where the reward is significantly higher ($r = 150$) and the defining constraint is not the lack of reviewers (we assume ample review capacity $K_{rev} = 1200$) but the extreme scarcity of publication slots ($K = 180$). This aligns with empirical findings on scientific gatekeeping, where high rejection rates are driven by physical slot limits rather than quality filtering alone (Siler et al., 2015; Casadevall and Fang, 2014), and it generates the stable coexistence topology that illustrates what we call the defence-by-congestion mechanism. By analysing both sets, we decouple the effects of review collapse from slot rationing, ensuring that our conclusions are robust across different segments of the academic publishing hierarchy.

6.2 The Decoupling of Capacities: Phase Diagram Analysis

Standard academic models often treat "journal capacity" as a single resource. Our dual-bottleneck framework reveals that the system's fate depends on how two distinct capacities interact: review capacity K_{rev}, which determines the screening intensity $\eta(x)$, and slot capacity K_{store}, which determines the rationing probability $\bar{\rho}(x)$. High K_{rev} prevents review collapse; low K_{store} creates what we call a defence by congestion.

To visualize the interaction between these forces, we perform a two-dimensional bifurcation scan across the parameter space (K_{rev}, K_{store}). The resulting phase diagram (Figure 8) maps the system's regimes.

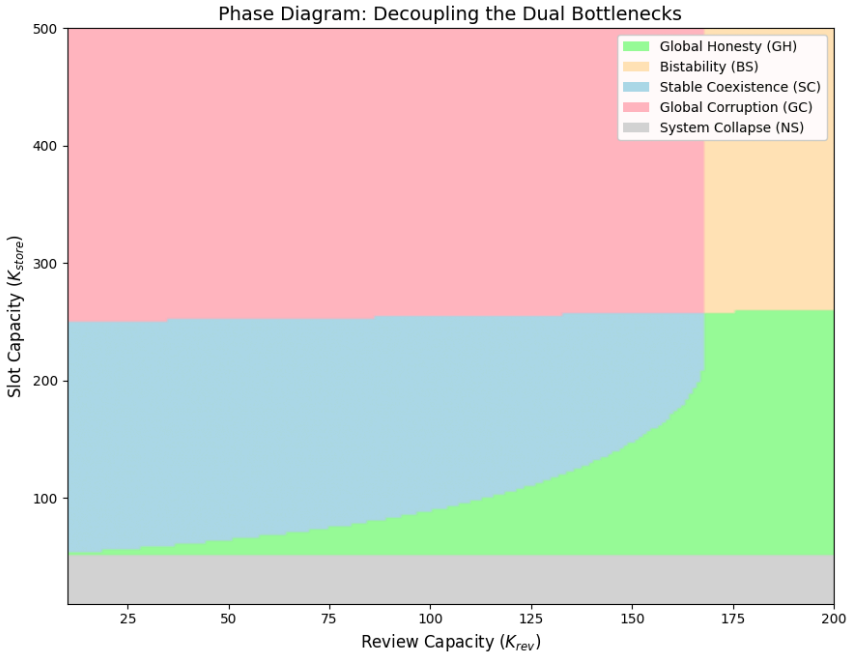

Figure 8: **Phase Diagram of the Dual-Bottleneck System.** The map classifies the topological regime based on the availability of review resources (K_{rev}) versus publication slots (K_{store}). **Green (Global Honesty):** High review capacity ensures cheating is unprofitable ($Q \to 1.0$). **Blue (Stable Coexistence):** Extreme slot scarcity ($K_{store} \ll N$) triggers the "Defense by Congestion," stabilizing the system at a mediocre but functional level. **Yellow (Bistability):** A transition zone where the outcome is history-dependent. **Red (Global Corruption):** When slots are abundant but review is weak, the defense mechanism vanishes, leading to total collapse ($Q \to \alpha$). **Gray (System Collapse):** Feasibility constraint violated (K_{store} too low); even honest authors cannot break even.

The topology of Figure 8 highlights two critical dynamical phenomena:

The paradox of slot expansion. Intuition suggests that giving a journal more publication slots should improve the system. We find the opposite. Moving vertically from the bottom of the diagram corresponds to increasing K_{store}. Initially, increasing slots rescues the system from total collapse (Gray to Blue), making participation rational for honest authors. However, further increasing K_{store} pushes the system from stable coexistence (Blue) into global corruption (Red) or bistability (Yellow). This paradox arises because tight storage constraints act as a passive defence by congestion, keeping $\bar{\rho}$ low enough to deter spam. By increasing K_{store}, the journal inadvertently dismantles this defence: $\bar{\rho}$ rises, the publication lottery becomes winnable again, and if K_{rev} is not increased simultaneously, the incentive to cheat reactivates and the system tips into full corruption. This finding aligns with strategic gatekeeping models showing that expanding capacity without scaling review rigour destroys the signal value of the journal (Muller-Itten, 2022).

The necessity of review resources. Moving horizontally to the right corresponds to investing in review capacity K_{rev}. Provided the system is in a feasible region (not Gray), we find that increasing K_{rev} is unconditionally beneficial, expanding the Global Honesty (Green) zone. Crucially, the boundary between corruption and honesty slopes upward: as K_{store} increases, the journal requires *disproportionately* more review capacity to maintain honesty. This implies that larger journals are inherently harder to police than smaller ones—they require a review infrastructure that scales super-linearly with volume to maintain the same level of integrity, exacerbating the already critical burden of peer review (Kovanis et al., 2016).

6.3 Population Pressure and Hysteresis: The Cost of Lost Trust

While capacity parameters define the structural limits of the journal, the total population N represents the external load. One might expect quality to decline smoothly as load increases. We find instead that the system exhibits hysteresis in the bistable regime: the value of N at which quality collapses is much higher than the value at which it recovers. This is consistent with theories of catastrophic shifts in complex systems, where state recovery requires significantly more effort than state maintenance (Scheffer et al., 2001).

To capture this path dependence, we use a numerical adiabatic continuation method. We define the system state $x_{st}(N)$ as the stable equilibrium selected by continuity from the previous state:

$$x_{st}(N_i) = \arg\min_{x \in \mathcal{E}_{st}(N_i)} |x - x_{st}(N_{i-1})| \tag{51}$$

where $\mathcal{E}_{st}(N_i)$ is the set of stable fixed points at N_i.

- **Forward Path (Crowding):** Increments N, initializing from $x_{st}(N_{min}) \approx 1$.
- **Backward Path (Recovery):** Decrements N, initializing from $x_{st}(N_{max}) \approx 0$.

Figure 9 illustrates the resulting trajectories using Baseline Set A.

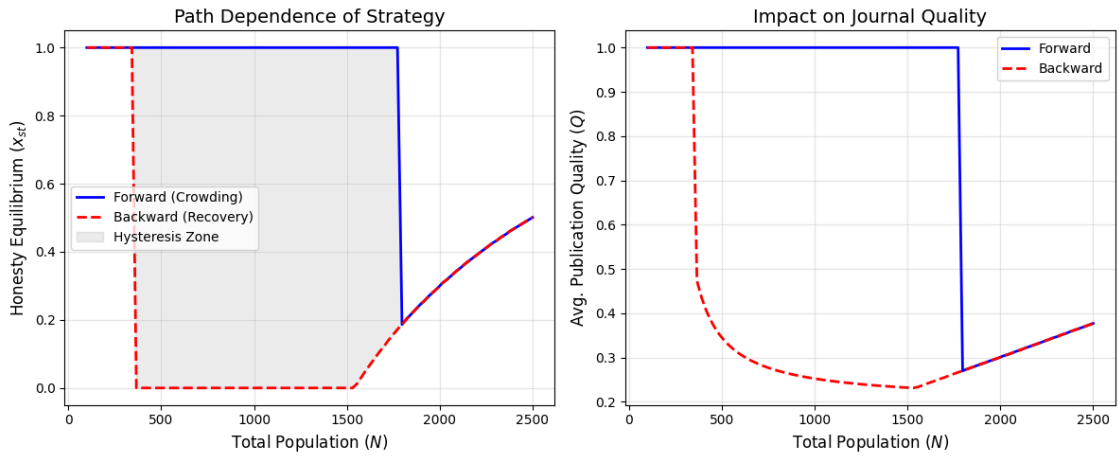

Figure 9: **Hysteresis Loop and Quality Collapse (Baseline Set A).** The shaded region represents the bistable interval $[N_\downarrow, N_\uparrow]$. **Left Panel (State):** On the Forward Path (Blue), the "Honest Norm" is resilient ($x_{st} \approx 1$) until the collapse threshold $N_\uparrow \approx 1800$. Once corrupted (Red Dashed), the system is trapped in the low-honesty attractor ($x_{st} \approx 0$). Recovery requires reducing N below the distinct recovery threshold $N_\downarrow \approx 350$. **Right Panel (Quality):** The collapse leads to a discrete jump in quality. In the corrupt state (Red Dashed), quality is dominated by the baseline productivity ($\alpha = 0.2$) but fluctuates slightly as review intensity $\eta(N)$ varies. Recovery only occurs when the system crosses $N_\downarrow$.

We highlight three features of the result.

Structural resilience (forward path). A journal with a strong reputation demonstrates remarkable resilience. As seen in the blue trajectory, the system maintains $x_{st} \approx 1$ and perfect quality ($Q = 1.0$) even when N far exceeds the nominal capacity, because honest authors self-regulate, keeping the effective load $V(1) = N\alpha$ manageable. The collapse point $N_\uparrow \approx 1800$ is the threshold where the honest equilibrium loses stability: $N_\uparrow = \inf\{N : \Pi_{bad}(1; N) \geq 0\}$.

The stickiness of corruption (backward path). Once the system settles into the corrupt equilibrium, the sheer volume of spam saturates the review capacity. Even if N is reduced, the system remains trapped in the low-quality state as long as $\Pi_{bad}(0; N) > 0$. The recovery point $N_\downarrow \approx 350$ corresponds to the threshold where the corrupt equilibrium finally becomes unstable. This asymmetry mirrors the dynamics of collective reputation, where a group remains locked in a low-quality equilibrium long after external conditions improve, because the burden of proof is too high for individual reformers (Tirole, 1996).

Crucially, we define the Average Publication Quality $Q(x)$ as the proportion of good papers in the passed pool $M(x)$:

$$Q(x) = \frac{N\alpha\pi_1(x)}{N\alpha\pi_1(x) + N(1-\alpha)(1-x)\pi_0(x)} \tag{52}$$

Since slot rationing ($\bar{\rho}$) applies uniformly to $M(x)$, it does not alter the ratio of good-to-bad papers. Consequently, quality remains low throughout the backward path until the critical transition occurs.

Policy implication: basin hopping. The disparity between $N_\uparrow$ and $N_\downarrow$ implies that incremental reforms are structurally insufficient for a corrupted journal. Within the bistable interval, marginal reductions in submissions do not change the long-run attractor. Escaping the low-quality basin requires interventions strong enough to move the system across the recovery threshold ($N < N_\downarrow$). This suggests that while soft measures work for maintenance, hard measures (e.g., temporary submission moratoriums) are structural necessities for recovery.

6.4 Strategic Levers and Regime Transitions

While capacity parameters define the structural limits, the strategic parameters $(r, c, \lambda, \epsilon, \alpha)$ determine the incentives. We perform a comparative sensitivity analysis across all five dimensions, grouping them into three functional categories: screening, economic incentives, and productivity.

Screening parameters: the false-positive rate versus friction. We begin with the two parameters controlling the peer review filter (Figure 10). Panel (a) reveals that the false-positive rate λ is the most critical parameter in the entire system. As λ increases, the system shifts rapidly from Green (Honesty) to Yellow (Bistability) and finally to Red (Corruption); in the yellow zone, a minor perturbation can trigger a catastrophic tip into the corrupt equilibrium. This aligns with signalling theory, where the reliability of the signal (low λ) is paramount for maintaining the separating equilibrium (Tiokhin et al., 2021).

In contrast, Panel (b) shows that friction ϵ has negligible structural impact: within the explored range, the regime boundaries remain largely unchanged. This implies that policies focused solely on reducing administrative friction for good papers yield little improvement in integrity compared to improving detection accuracy.

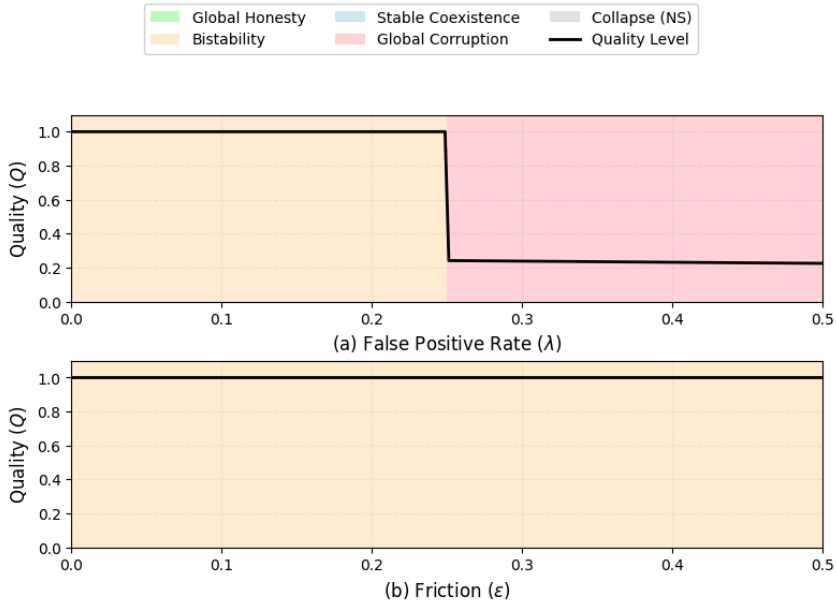

Figure 10: **Screening Parameters and Regime Stability. (a) False Positive Rate λ:** Acts as a major destabilizer. Increasing λ destroys the honest equilibrium, pushing the system into high-risk bistability. **(b) Friction ϵ:** Has negligible impact on regime boundaries, acting merely as a silent administrative cost.

Economic levers: reward and submission cost. We next examine the financial and reputational incentives (Figure 11). Panel (c) demonstrates that high rewards r act as an accelerant for corruption, significantly expanding the Red and Yellow zones. When publication rewards are excessive, the lottery of rubber-stamping becomes profitable even at low probabilities—confirming the mechanism of "natural selection of bad science," where high incentives select for strategies that maximize output at the expense of quality (Smaldino and McElreath, 2016).

While reward drives the system toward corruption, submission cost c pulls it back. Panel (d) shows that increasing c effectively shifts the system from Red to Green, providing theoretical support for submission fees as a tool to reduce congestion and filter out low-quality submissions (Cotton, 2013). However, pushing costs too high forces the system into the Gray zone, where even honest authors cannot sustain participation ($\Pi_{good} < 0$).

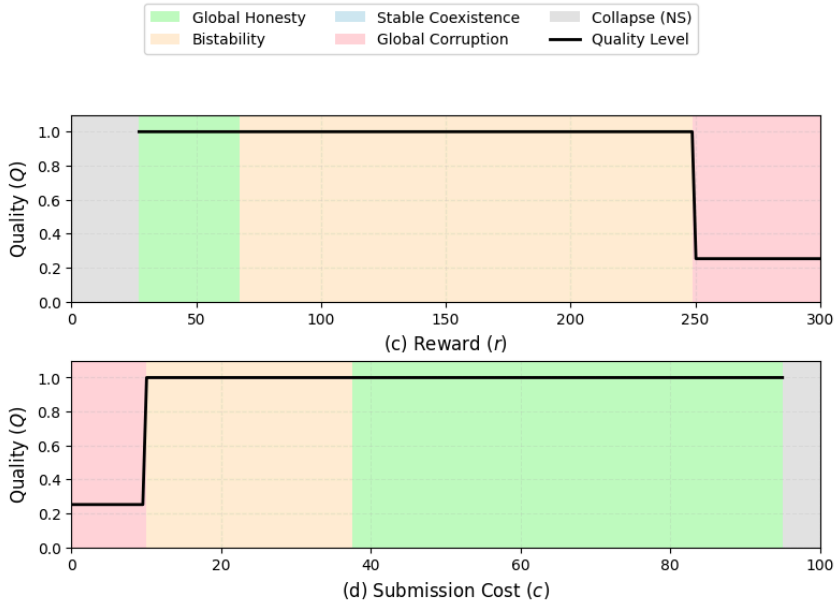

Figure 11: **Economic Incentives and Filtering. (c) Reward r:** High rewards expand the corrupt zones (Red/Yellow), incentivizing spamming strategies. **(d) Cost c:** Submission fees filter out spam, restoring Green, but excessive costs kill the journal (Gray).

The productivity paradox. Finally, we analyse the impact of aggregate author productivity (Figure 12). One might expect that higher productivity α would stabilize the system, since the input mix improves. We find the opposite: as α increases, the background topology shifts toward unstable regimes (Yellow/Red), even though the average quality (black line) rises mechanically. The mechanism is congestion. As legitimate authors produce more papers, the total volume rises and consumes scarce review capacity; if K_{rev} is fixed, this "good congestion" dilutes the screening intensity for everyone, paradoxically making the system more vulnerable to invasion by spammers.

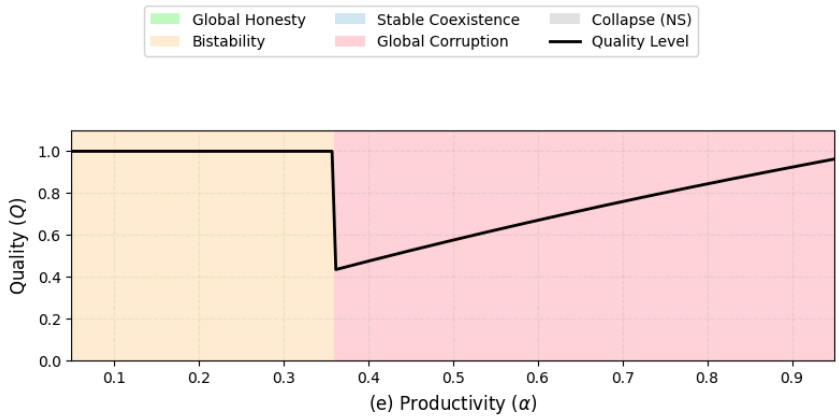

Figure 12: **The Productivity Paradox. (e) Productivity α:** Higher productivity increases total congestion. While observable quality rises (black line), the structural stability degrades (background shifts to Red), narrowing the window of stability.

7 Agent-Based Modeling

7.1 Simulation Design

The analytical results above assume a large, well-mixed population and deterministic strategy evolution—the standard mean-field approximation (Hofbauer and Sigmund, 1998). Real academic communities are finite and subject to stochastic fluctuations (Nowak, 2006). We therefore implement an agent-based model (ABM) to test whether our main claims survive these complications: that dominated strategies die out naturally, that the predicted long-run outcomes emerge from local interactions under noise, and that the defence-by-congestion mechanism is not an artefact of the mean-field limit (Bonabeau, 2002).

We simulate a population of $N = 1000$ agents, following the standard protocols of stochastic evolutionary game theory in finite populations Traulsen and Hauert (2009). We initialise the population with a uniform distribution across all four strategies (25% each), so that the simulation must find the reduced state on its own rather than having it assumed. The process proceeds as follows:

1. **Interaction Phase:** In each time step t, agents generate submissions based on their full strategy space ($s_i \in \{\text{OG}, \text{AS}, \text{OB}, \text{NS}\}$). The aggregate volume V and peer-accepted volume M are calculated endogenously, determining the congestion levels η and ρ.

2. **Payoff Realization:** Agents receive payoffs Π_i based on the realized publication outcomes. This captures the "lottery" nature of the congestion bottleneck.

3. **Strategy Update (Social Learning):** We employ the **pairwise comparison process** to model strategy evolution. A randomly selected agent A considers mimicking agent B's strategy. The

probability of switching $s_A \to s_B$ is given by the Fermi function derived from statistical mechanics Szabó and Tőke (1998):

$$P(s_A \to s_B) = \frac{1}{1 + \exp(-\beta[\Pi_B - \Pi_A])} \tag{53}$$

where β represents the intensity of selection. This update rule is the standard method for introducing bounded rationality and decision noise into evolutionary dynamics Traulsen and Hauert (2009).

7.2 Simulation Results

We simulated both parameter sets from Table 5 with a high selection intensity ($\beta = 1.0$) to reduce noise. Figure 13 shows that the agent-based simulations corroborate both the strategy elimination and the macroscopic predictions from Section 3.

In both scenarios, the "Only Bad" (OB) and "Never Submit" (NS) strategies are rapidly driven to extinction within the first few hundred time steps, directly verifying the iterated elimination of strictly dominated strategies (IESDS) conducted in Section 2.6. The system autonomously reduces to a binary contest between OG and AS, consistent with the finding that strictly dominated strategies cannot survive under monotonic selection dynamics Samuelson (1997).

In the review-dominated regime (Scenario A), the dynamics exhibit sharp bistability (Figure 13a). Because the initial proportion of honest OG authors (25%) falls well below the theoretical unstable tipping point ($x^* \approx 0.82$), the system cannot self-correct. Once OB and NS vanish, the AS strategy outcompetes the remaining honest authors, driving the system into the absorbing state of full corruption ($x = 0$). This confirms the severe path dependence predicted by the analytical model Scheffer et al. (2001); Schelling (1978).

In the rationing-dominated regime (Scenario B), we observe a qualitatively different outcome (Figure 13b). Despite the initial chaotic interaction of all four strategies, the population rapidly settles to a long-run average of $x^* \approx 0.49$ for OG and AS, while OB and NS remain extinct. This result aligns precisely with the analytical fixed point ($x^* \approx 0.49$) identified in the phase portrait analysis, confirming that the defence-by-congestion mechanism is a robust emergent property, immune to initial strategy turbulence and finite-size stochasticity.

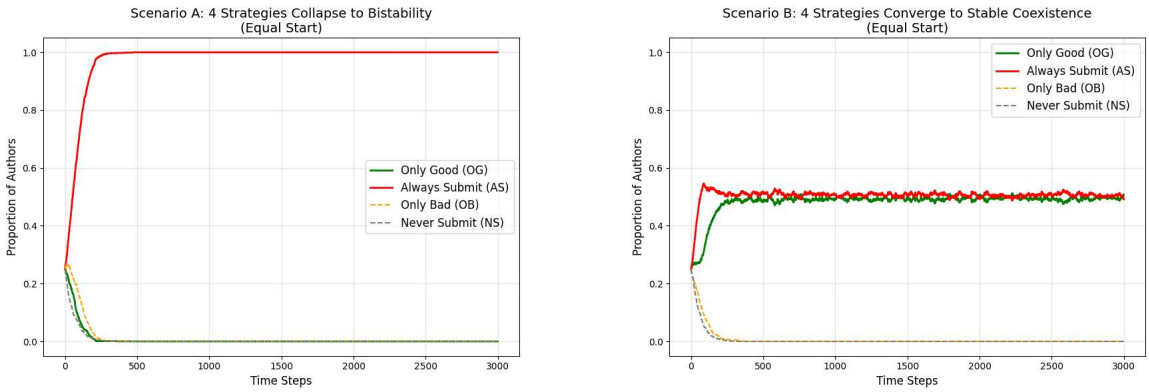

(a) Scenario A: Review-Dominated Dynamics

(b) Scenario B: Stable Coexistence

Figure 13: **Agent-Based Model Trajectories with 4 Strategies ($\beta = 1.0$).** The plots show the time evolution of the strategy proportions. **(a)** Strong Bistability: Starting from an equal uniform mix, the system rapidly eliminates OB and NS. Because the initial OG fraction (0.25) is below the critical tipping point, the system collapses to full corruption ($x = 0$). **(b)** Stable Coexistence: OB and NS are quickly eliminated, and the remaining OG and AS populations converge to a mixed equilibrium ($x^* \approx 0.49$). The precise agreement with analytical predictions confirms both the theoretical dominance analysis and the structural congestion stabilization.

7.3 Co-operation ABM

To rigorously validate the robustness of our theoretical findings and explore the microscopic foundations of the "Defense by Congestion" mechanism, we develop a specialized Co-operation Agent-Based Model (ABM) featuring local batch-matching. While the standard replicator dynamics rely on a mean-field approximation assuming an infinite population and a single monolithic market, real-world academic publishing is inherently fragmented. Authors typically compete within specific sub-disciplines, specialized journals, or distinct temporal submission windows. The primary motivation for designing this ABM is to investigate whether the macroscopic stability predicted by the ODEs holds when subjected to the localized uncertainty, heterogeneous payoffs, and stochastic fluctuations characteristic of finite populations.

The core innovation of this extended ABM is the *batch-matching* algorithm. Instead of evaluating all N agents globally, the population is randomly partitioned into local batches (sub-markets) in each time step. The journal's review and storage capacities (K_{rev} and K_{store}) are scaled proportionally to match these local batches. This approach significantly improves upon standard global ABMs by introducing critical variance into the localized rationing probability (ρ) and enabling stochastic drift. It captures the granular reality that the publication "lottery" is not uniform; different sub-fields experience varying degrees of congestion simultaneously. This localized competition provides a much more rigorous stress-test for the theoretical equilibria.

The simulation results validate our threshold competition framework across two key dimensions, as

illustrated in Figure 14:

Trajectory Analysis and Robustness: Figures 14(a) and 14(b) illustrate the time evolution of strategies under batch matching. In the rationing-dominated regime (Scenario B), despite chaotic local interactions, both the theoretical ODE and the ABM trajectories converge tightly to the stable internal equilibrium of $x^* \approx 0.49$. This confirms that the stable coexistence induced by congestion is highly robust against local noise. Conversely, in the review-dominated regime (Scenario A), the ABM demonstrates sharp bistability, collapsing rapidly to global honesty ($x = 1.0$) or global corruption ($x = 0.0$) depending on initial conditions. The ABM closely traces the expected theoretical paths but exhibits micro-oscillations, validating the IESDS reduction while underscoring the severe path-dependence of the bistable regime.

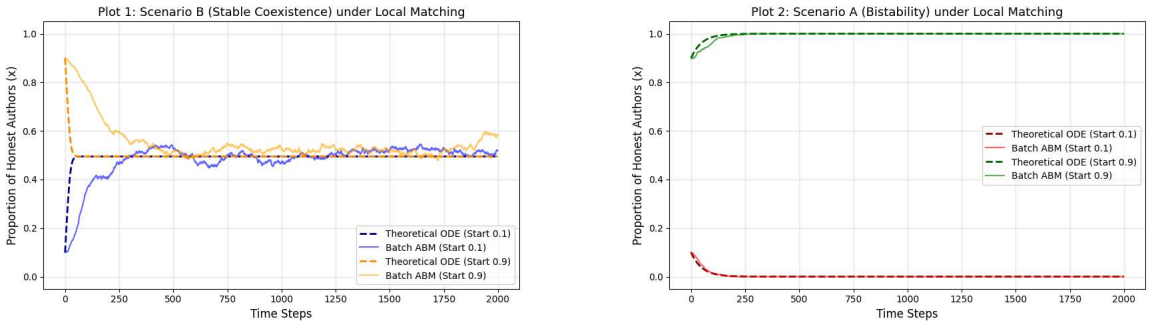

(a) Scenario B (Stable Coexistence) Trajectories

(b) Scenario A (Bistability) Trajectories

Figure 14: **Results of the Co-operation ABM with Local Batch-Matching.** The composite figure demonstrates that the theoretical mean-field predictions hold even under the rigorous localized constraints and stochastic noise introduced by fragmented sub-markets.

The Journal as an Independent Evolutionary Agent

From Fixed Environment to Co-Evolutionary Agent In previous sections, including the local batch-matching analysis in Section 5.4, the journal was mathematically treated as a fixed institutional player locked into the "Only Good" (OG) strategy. It acted purely as a static environmental constraint enforcing a peer review filter and a capacity limit. By elevating the journal to an independent, adaptive agent, we transition the framework from a simple population game into an asymmetric co-evolutionary game. Journals now possess their own complete strategy space, consisting of Only Good (OG), Always Accept (AA), Only Bad (OB), and Reject All (AR). They update their behavior based on their own distinct utility function, which must balance the reputation benefits of publishing good papers, the penalties of publishing bad papers, and the administrative costs of peer review. This structural update captures the "Red Queen" dynamic inherent in academic publishing: if the author population becomes overwhelmingly honest, journals are financially incentivized to drop their costly screening processes and

adopt the "Always Accept" (AA) strategy. However, this relaxation immediately triggers an opportunistic invasion of "Always Submit" (AS) authors, which in turn forces the journals to revert to strict filtering.

The Significance of the 10:1 Local Batch-Matching Mechanism Real-world academic publishing is not a monolithic global market, but rather a highly fragmented ecosystem divided into specific sub-disciplines and specialized journals. To accurately model this micro-level reality, we introduce an extreme 10:1 local batch-matching mechanism, pairing 1000 authors with 100 distinct journal agents. In this configuration, global capacities (K_{rev} and K_{store}) are scaled down proportionally to match the size of these local sub-markets. The primary significance of this 10:1 ratio is the maximization of localized variance and stochastic drift. By breaking the assumption of uniform global congestion, a single journal agent shifting to a rubber-stamping (AA) strategy creates a massive, localized payoff spike for the specific batch of 10 authors assigned to it, severely testing the structural resilience of the system.

ABM Configuration and Result Analysis To simulate this environment, we employ an Agent-Based Model (ABM) where both populations update their strategies via a Fermi pairwise comparison process, introducing bounded rationality and decision noise.

First, we analyze the global co-evolutionary baseline without fragmentation, as illustrated in **Figure 15**. The ABM realization demonstrates that structurally dominated strategies, such as Only Bad (OB) and Reject All (AR), are rapidly driven to extinction. The remaining strategies enter a dynamic coexistence where the continuous mathematical predictions (ODE) and the finite-population ABM trajectories align closely.

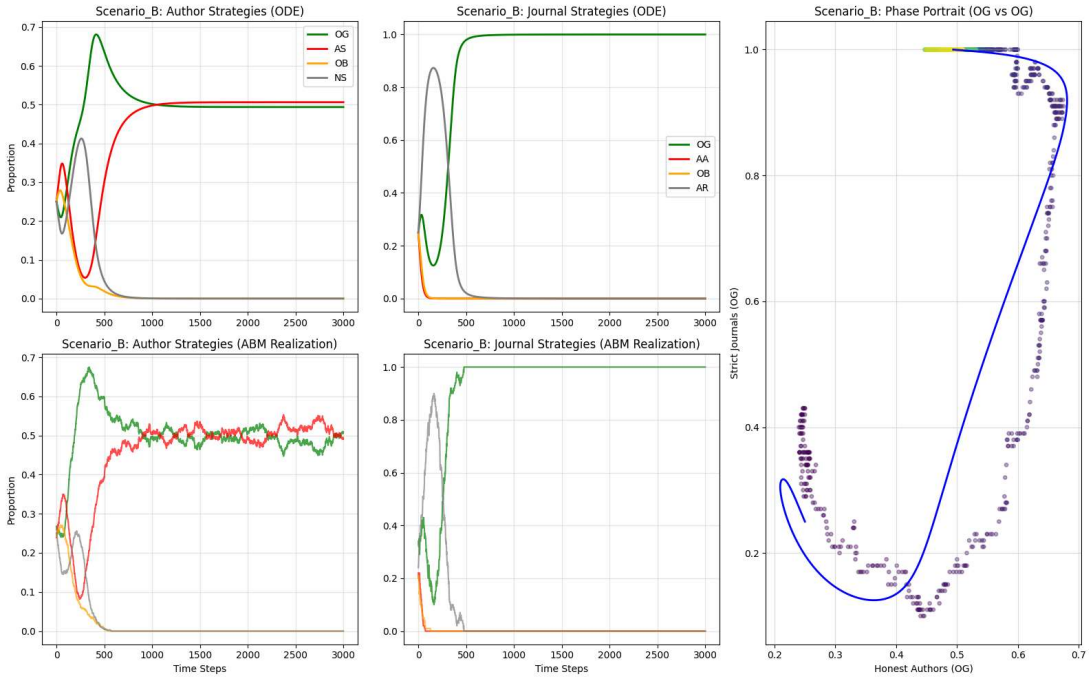

Figure 15: Global Co-evolutionary Baseline (Scenario B). The system demonstrates stable coexistence without fragmentation, where the finite-population ABM trajectory closely aligns with the continuous ODE predictions.

However, under the extreme fragmentation of the 10:1 batch-matching architecture, illustrated in **Figure 16**, the system exhibits a profound deviation from the global mean-field predictions. In the rationing-dominated regime (Scenario B), the ABM reveals that the "Never Submit" (NS) strategy unexpectedly overtakes the ecosystem. This occurs due to a phenomenon we term "Feasibility Collapse." Because the global publication slots (K_{store}) are strictly divided among the 100 journals, each local sub-market operates with an extremely narrow bottleneck. If even a small number of opportunistic (AS) authors flood a specific 10-person room, the local rationing probability (ρ) plummets below the economic break-even threshold required to justify the submission cost (c). Consequently, the "Defense by Congestion" becomes too punitive; even honest authors (OG) suffer negative expected payoffs and are forced to exit the market entirely, leaving NS as the only rational, risk-averse strategy.

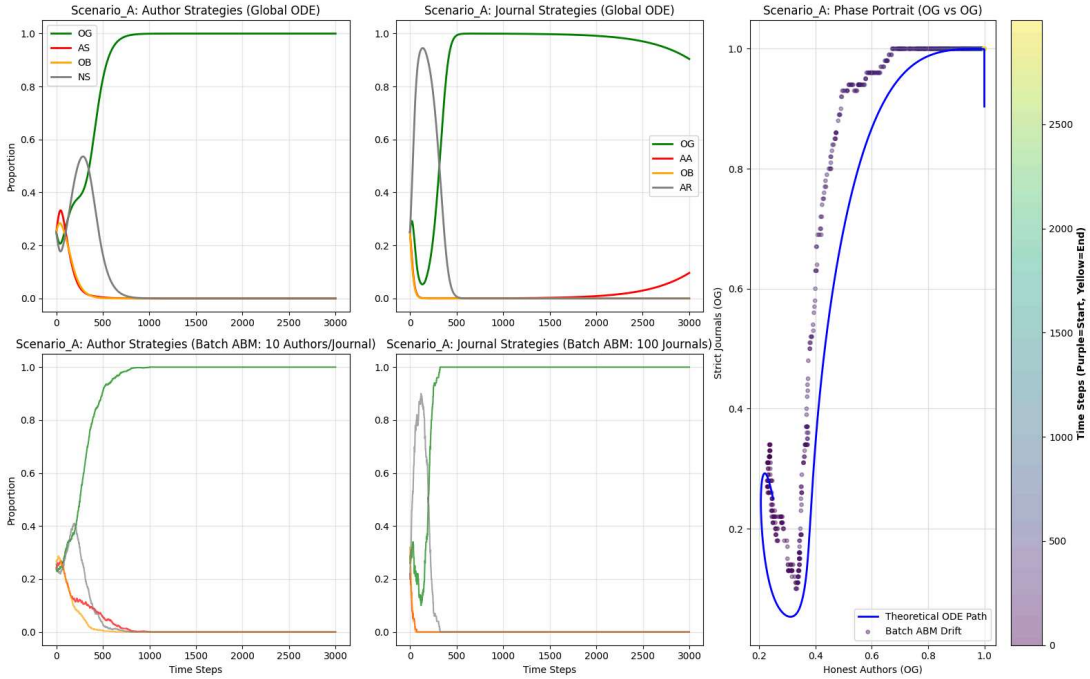

Figure 16: Co-evolutionary Dynamics under 10:1 Local Batch-Matching (Scenario B). The extreme fragmentation induces severe localized congestion variance, causing the "Never Submit" (NS) strategy to dominate due to Feasibility Collapse.

8 Conclusion

References

Balazs Aczél, Barnabas Szaszi, and Alex O. Holcombe. A billion-dollar donation: estimating the cost of researchers' time spent on peer review. *Research Integrity and Peer Review*, 6(14), 2021. doi: 10.1186/s41073-021-00118-2.

George A Akerlof. The market for "lemons": Quality uncertainty and the market mechanism. *The Quarterly Journal of Economics*, 84(3):488–500, 1970. doi: 10.2307/1879431.

arXiv. arXiv sets new record for monthly submissions. arXiv Blog, 2024. URL `https://blog.arxiv.org/2024/11/04/arxiv-sets-new-record-for-monthly-submissions-again/`. Accessed March 2026.

Ofer H Azar. The review process in economics: Is it too fast? *Southern Economic Journal*, 72(2): 482–491, 2005.

Federico Bianchi, Francisco Grimaldo, Giangiacomo Bravo, and Flaminio Squazzoni. The peer review game: an agent-based model of scientists facing resource constraints and institutional pressures. *Scientometrics*, 116(3):1401–1420, 2018. doi: 10.1007/s11192-018-2825-4.

Eric Bonabeau. Agent-based modeling: Methods and techniques for simulating human systems. *Proceedings of the National Academy of Sciences*, 99(suppl_3):7280–7287, 2002.

Lutz Bornmann. Scientific peer review. *Annual Review of Information Science and Technology*, 45(1): 197–245, 2011.

Lutz Bornmann, Rüdiger Mutz, and Hans-Dieter Daniel. A reliability-generalization study of journal peer reviews: A multilevel meta-analysis of inter-rater reliability and its determinants. *PLoS One*, 5 (12):e14331, 2010.

Arturo Casadevall and Ferric C Fang. Reforming science: methodological and cultural reforms. *Infection and Immunity*, 82(12):4881–4888, 2014.

Alessandro Checco, Lorenzo Bracciale, Pierpaolo Loreti, Stephen Pinfield, and Giuseppe Bianchi. AI-assisted peer review. *Humanities and Social Sciences Communications*, 8(1):25, 2021. doi: 10.1057/s41599-020-00703-8.

Domenic V Cicchetti. The reliability of peer review for manuscript and grant submissions: A cross-disciplinary investigation. *Behavioral and Brain Sciences*, 14(1):119–135, 1991.

Christopher Cotton. Submission fees and response times in academic publishing. *American Economic Review*, 103(1):501–509, 2013.

Glenn Ellison. Evolving standards for academic publishing: A q-r theory. *Journal of Political Economy*, 110(5):994–1034, 2002.

Drew Fudenberg and Jean Tirole. *Game Theory*. MIT Press, Cambridge, MA, 1991.

Paul Ginsparg. arXiv at 20. *Nature*, 476(7359):145–147, 2011. doi: 10.1038/476145a.

Kevin Gross and Carl T Bergstrom. Contest models highlight inherent inefficiencies of scientific funding competitions. *PLOS Biology*, 17(1):e3000065, 2019. doi: 10.1371/journal.pbio.3000065.

Josef Hofbauer and Karl Sigmund. *Evolutionary Games and Population Dynamics*. Cambridge University Press, Cambridge, UK, 1998.

John P A Ioannidis. Why most published research findings are false. *PLOS Medicine*, 2(8):e124, 2005. doi: 10.1371/journal.pmed.0020124.

Dmitry Kobak, Rita González-Márquez, Emőke-Ágnes Horvát, and Jan Lause. Delving into LLM-assisted writing in biomedical publications through excess vocabulary. *Science Advances*, 11(27), 2025. doi: 10.1126/sciadv.adt3813.

Michail Kovanis, Raphael Porcher, Philippe Ravaud, and Ludovic Trinquart. The global burden of journal peer review in the biomedical literature: Strong imbalance in the collective enterprise. *PLoS One*, 11(11):e0166387, 2016.

Giuseppe Russo Latona, Manoel Horta Ribeiro, Tim R. Davidson, Veniamin Veselovsky, and Robert West. The AI review lottery: Widespread AI-assisted peer reviews boost paper scores and acceptance rates, 2024.

Weixin Liang, Mert Yuksekgonul, Yining Mao, Eric Wu, and James Zou. GPT detectors are biased against non-native English writers. *Patterns*, 4(7):100779, 2023. doi: 10.1016/j.patter.2023.100779.

Weixin Liang, Zachary Izzo, Yaohui Zhang, Haley Lepp, Hancheng Cao, Xuandong Zhao, Lingjiao Chen, Haoyu Ye, Sheng Liu, Zhi Huang, Daniel A. McFarland, and James Y. Zou. Monitoring AI-modified content at scale: A case study on the impact of ChatGPT on AI conference peer reviews. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 29575–29620. PMLR, 2024. URL `https://proceedings.mlr.press/v235/liang24b.html`.

Weixin Liang, Yaohui Zhang, Zhengxuan Wu, Haley Lepp, Hancheng Cao, Xuandong Zhao, Lingjiao Chen, Haoyu Ye, Sheng Liu, Zhi Huang, Daniel A. McFarland, and James Y. Zou. Quantifying large language model usage in scientific papers. *Nature Human Behaviour*, 9:2599–2609, 2025. doi: 10.1038/s41562-025-02273-8.

Laura A. Meyerson, Adam Suzzi-Simmons, and Daniel Simberloff. Quantifying reviewer declines in scientific publishing: twenty-one years of data from Biological Invasions 2002–2024. *Biological Invasions*, 27:223, 2025. doi: 10.1007/s10530-025-03679-1.

Michele Muller-Itten. Gatekeeper competition. *European Economic Review*, 148:104230, 2022.

Miryam Naddaf. More than half of researchers now use AI for peer review — often against guidance. *Nature*, 649, 2025. doi: 10.1038/d41586-025-04066-5.

Martin A Nowak. *Evolutionary dynamics: exploring the equations of life*. Harvard university press, 2006.

Shakked Noy and Whitney Zhang. Experimental evidence on the productivity effects of generative artificial intelligence. *Science*, 381(6654):187–192, 2023. doi: 10.1126/science.adh2586.

Robert W Rosenthal. A class of games possessing pure-strategy nash equilibria. *International Journal of Game Theory*, 2(1):65–67, 1973.

Larry Samuelson. *Evolutionary games and equilibrium selection*. MIT press, 1997.

William H Sandholm. *Population games and evolutionary dynamics*. MIT press, Cambridge, MA, 2010.

Marten Scheffer, Steve Carpenter, Jonathan A Foley, Carl Folke, and Brian Walker. Catastrophic shifts in ecosystems. *Nature*, 413(6856):591–596, 2001. doi: 10.1038/35098000.

Thomas C Schelling. *Micromotives and Macrobehavior*. W. W. Norton & Company, New York, 1978. ISBN 9780393090090.

Kyle Siler, Kirby Lee, and Lisa Bero. Measuring the effectiveness of scientific gatekeeping. *Proceedings of the National Academy of Sciences*, 112(2):360–365, 2015.

Paul E Smaldino and Richard McElreath. The natural selection of bad science. *Royal Society Open Science*, 3(9):160384, 2016. doi: 10.1098/rsos.160384.

Paula Stephan. *How Economics Shapes Science*. Harvard University Press, Cambridge, MA, 2012.

Chris Stokel-Walker and Richard Van Noorden. What ChatGPT and generative AI mean for science. *Nature*, 614(7947):214–216, 2023. doi: 10.1038/d41586-023-00340-6.

György Szabó and Csaba Tőke. Evolutionary prisoner's dilemma game on a square lattice. *Physical Review E*, 58(1):69, 1998.

Stefan Thurner and Rudolf Hanel. Peer-review in a world with rational scientists: Toward a selection of the average. *The European Physical Journal B*, 84(4):707–711, 2011.

Leonid Tiokhin, Karthik Panchanathan, Daniel Lakens, Simine Vazire, Thomas Morgan, and Kevin Zollman. Honest signaling in academic publishing. *PLoS One*, 16(2):e0246675, 2021.

Jean Tirole. A theory of collective reputations (with applications to the persistence of corruption and to firm quality). *The Review of Economic Studies*, 63(1):1–22, 1996. doi: 10.2307/2298112.

Arne Traulsen and Christoph Hauert. Stochastic evolutionary game dynamics. *Reviews of nonlinear dynamics and complexity*, 2:25–61, 2009.

Jörgen W Weibull. *Evolutionary game theory.* MIT press, Cambridge, MA, 1995.

Kevin J. S. Zollman, Julian García, and Toby Handfield. Academic journals, incentives, and the quality of peer review: A model. *Philosophy of Science*, 91(1):186–203, 2024.